


```

In [1]: # Fresh vs Rotten Fruit Classification - Multi-Task Learning
# Complete Conference-Ready Implementation with Enhanced Preprocessing and Improved ViT
# FINAL VERSION WITH ALL ISSUES FIXED FOR 9 MODELS
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
import os
import warnings
warnings.filterwarnings('ignore')

# Deep Learning Libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, models, optimizers, regularizers
from tensorflow.keras.applications import VGG16, InceptionV3, ResNet152V2
from tensorflow.keras.applications import EfficientNetB4, DenseNet121
from tensorflow.keras.applications import ConvNeXtBase
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
from sklearn.utils.class_weight import compute_class_weight

# FIXED: Add tensorflow_hub import for ViT model
import tensorflow_hub as hub

# Metrics and Evaluation
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

# Set style for publication-quality figures
plt.style.use('default')
sns.set_palette("husl")
plt.rcParams.update({
    'font.size': 12,
    'axes.titlesize': 14,
    'axes.labelsize': 12,
    'xtick.labelsize': 10,
    'ytick.labelsize': 10,
    'legend.fontsize': 10,
    'figure.titlesize': 16
})

# Set random seeds for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# GPU configuration
gpus = tf.config.experimental.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:

```

```

        tf.config.experimental.set_memory_growth(gpu, True)
        print(f"GPU available: {len(gpus)} GPU(s)")
    except RuntimeError as e:
        print(e)
else:
    print("GPU not available, using CPU")

print(f"TensorFlow version: {tf.__version__}")
print(f"Keras version: {keras.__version__}")

#=====
# ENHANCED DATA LOADING AND PREPROCESSING - FIXED VERSION
#=====

def normalize_fruit_name(folder_name):
    """
    Enhanced fruit name normalization with comprehensive spelling fixes
    """
    folder_lower = folder_name.lower().strip()

    # Remove fresh/rotten prefixes
    if folder_lower.startswith('fresh'):
        fruit_name = folder_lower[5:].strip()
        freshness = 1 # Fresh = 1
    elif folder_lower.startswith('rotten'):
        fruit_name = folder_lower[6:].strip()
        freshness = 0 # Rotten = 0
    else:
        fruit_name = folder_lower
        # Try to infer freshness from name
        if 'rotten' in fruit_name or 'bad' in fruit_name:
            freshness = 0
        else:
            freshness = 1 # Default to fresh

    # Comprehensive spelling corrections and normalization
    corrections = {
        'tamto': 'tomato',
        'patato': 'potato',
        'bittergroud': 'bittergourd',
        'apsicum': 'capsicum',
        'apples': 'apple',
        'oranges': 'orange',
        'bananas': 'banana',
        'tomatos': 'tomato',
        'potatos': 'potato',
        'cucumbers': 'cucumber'
    }

    fruit_name = corrections.get(fruit_name, fruit_name)

    return fruit_name, freshness

def create_dummy_data():
    """Create dummy data for testing when actual dataset is not available"""
    print("⚠ Creating dummy dataset for testing...")
    fruits = ['apple', 'banana', 'orange', 'tomato', 'potato', 'capsicum', 'cu

```

```

cumber', 'okra', 'bittergourd']
dummy_data = []

for fruit in fruits:
    for fresh in [0, 1]:
        for i in range(200): # Increased from 150 to 200
            dummy_data.append({
                'filepath': f'dummy_{fruit}_{fresh}_{i}.jpg',
                'fruit': fruit,
                'fresh': fresh
            })

df = pd.DataFrame(dummy_data)
print(f"🔧 Dummy dataset created with {len(df)} samples")
return df

def load_and_process_dataset(base_path='/kaggle/input/fresh-and-stale-classifi-
cation'):
    """
    Enhanced dataset loading with comprehensive preprocessing pipeline - FIXED
    VERSION
    """
    print("🔄 Starting Enhanced Dataset Loading and Processing...")

    TRAIN_PATH = os.path.join(base_path, 'dataset/Train')
    TEST_PATH = os.path.join(base_path, 'dataset/Test')

    def load_data_from_path(path):
        """Load data from a specific path"""
        filepaths, fruits, freshness = [], [], []

        if not os.path.exists(path):
            print(f"⚠️ Path {path} does not exist")
            return pd.DataFrame({'filepath': [], 'fruit': [], 'fresh': []})

        try:
            folders = [f for f in os.listdir(path) if os.path.isdir(os.path.jo-
in(path, f))]
        except Exception as e:
            print(f"⚠️ Error reading directory {path}: {e}")
            return pd.DataFrame({'filepath': [], 'fruit': [], 'fresh': []})

        for folder in tqdm(folders, desc=f"Processing {os.path.basename(pat-
h)}"):
            try:
                folder_path = os.path.join(path, folder)
                fruit_name, fresh_value = normalize_fruit_name(folder)

                # Get all image files
                image_files = [f for f in os.listdir(folder_path)
                               if f.lower().endswith(('.png', '.jpg', '.jpeg',
'.bmp', '.tiff'))]

                for img_file in image_files:
                    full_path = os.path.join(folder_path, img_file)
                    # Only add if file actually exists
                    if os.path.exists(full_path):

```

```

        filepaths.append(full_path)
        fruits.append(fruit_name)
        freshness.append(fresh_value)
    except Exception as e:
        print(f"⚠️ Error processing folder {folder}: {e}")
        continue

    return pd.DataFrame({
        'filepath': filepaths,
        'fruit': fruits,
        'fresh': freshness
    })

# Load train and test data
df_train = load_data_from_path(TRAIN_PATH)
df_test = load_data_from_path(TEST_PATH)

# Combine datasets
df_all = pd.concat([df_train, df_test], axis=0, ignore_index=True)

if len(df_all) == 0:
    print("⚠️ No data found, creating dummy data...")
    return create_dummy_data()

print(f"✅ Total images loaded: {len(df_all)}")
print(f"🍎 Fruit distribution:")
fruit_counts = df_all['fruit'].value_counts()
print(fruit_counts)

# Target 9 fruit classes
target_fruits = ['apple', 'banana', 'capsicum', 'cucumber', 'okra', 'orange', 'potato', 'tomato', 'bittergourd']

# Filter to target fruits only
df_filtered = df_all[df_all['fruit'].isin(target_fruits)].copy()

if len(df_filtered) == 0:
    print("⚠️ No target fruits found, creating dummy data...")
    return create_dummy_data()

print(f"\n📄 Filtered to target fruits: {len(df_filtered)} images")
print(f"🎯 Target fruit distribution:")
filtered_counts = df_filtered['fruit'].value_counts()
print(filtered_counts)

return df_filtered

def apply_smart_balancing(df):
    """
    Apply smart balancing strategy - FIXED VERSION
    """
    print(f"\n🧠 Applying Smart Balancing Strategy...")

    if len(df) == 0:
        print("⚠️ Empty dataframe received, creating dummy data...")
        return create_dummy_data()

```

```

balanced_data = []
augmentation_stats = {}

for fruit in df['fruit'].unique():
    fruit_data = df[df['fruit'] == fruit].copy()
    total_samples = len(fruit_data)

    print(f"\n🍎 Processing {fruit}: {total_samples} samples")

    if total_samples >= 800: # Reduced threshold from 1000 to 800
        # Case 1: Sufficient data, sample up to 3000
        target_size = min(3000, total_samples)

        # Stratified sampling to maintain fresh/rotten balance
        fresh_data = fruit_data[fruit_data['fresh'] == 1]
        rotten_data = fruit_data[fruit_data['fresh'] == 0]

        if len(fresh_data) == 0:
            fresh_ratio = 0
        elif len(rotten_data) == 0:
            fresh_ratio = 1
        else:
            fresh_ratio = len(fresh_data) / total_samples

        target_fresh = max(1, int(target_size * fresh_ratio))
        target_rotten = target_size - target_fresh

        sampled_fresh = fresh_data.sample(n=min(target_fresh, len(fresh_data)),
                                           random_state=42, replace=len(fresh_data) < target_fresh)
        sampled_rotten = rotten_data.sample(n=min(target_rotten, len(rotten_data)),
                                           random_state=42, replace=len(rotten_data) < target_rotten)

        fruit_balanced = pd.concat([sampled_fresh, sampled_rotten])
        print(f" ✅ Sampled {len(fruit_balanced)} images (Fresh: {len(sampled_fresh)}, Rotten: {len(sampled_rotten)})")

    else:
        # Case 2: Insufficient data, augment to reach minimum threshold
        target_size = max(800, total_samples) # Minimum 800 samples

        if total_samples < 100:
            print(f" ⚠️ WARNING: {fruit} has only {total_samples} original images (<100)")

        # Calculate how many augmented images we need
        needed_samples = target_size - total_samples

        if needed_samples > 0:
            print(f" 🔄 Need to augment {needed_samples} additional images")

            # Create augmented entries (not actual images, just metadata)
            augmented_data = []

```

```

samples_per_original = max(1, needed_samples // total_samples)

for _, row in fruit_data.iterrows():
    for i in range(samples_per_original):
        aug_row = row.copy()
        # Mark as augmented with unique identifier
        aug_row['filepath'] = f"{row['filepath']}_aug_{i}"
        aug_row['is_augmented'] = True
        augmented_data.append(aug_row)

# Add remaining samples if needed
remaining = needed_samples - len(augmented_data)
if remaining > 0:
    for i in range(remaining):
        orig_idx = i % len(fruit_data)
        aug_row = fruit_data.iloc[orig_idx].copy()
        aug_row['filepath'] = f"{aug_row['filepath']}_aug_extra_{i}"
        aug_row['is_augmented'] = True
        augmented_data.append(aug_row)

# Mark original data
fruit_data_marked = fruit_data.copy()
fruit_data_marked['is_augmented'] = False

# Combine original and augmented
fruit_balanced = pd.concat([fruit_data_marked, pd.DataFrame(augmented_data)]).reset_index(drop=True)
print(f"✅ Created {len(fruit_balanced)} samples ({total_samples} original + {len(augmented_data)} augmented)")

augmentation_stats[fruit] = {
    'original': total_samples,
    'augmented': len(augmented_data),
    'total': len(fruit_balanced)
}
else:
    fruit_data['is_augmented'] = False
    fruit_balanced = fruit_data
    print(f"✅ Using {len(fruit_balanced)} original samples")

balanced_data.append(fruit_balanced)

# Combine all balanced data
df_balanced = pd.concat(balanced_data, ignore_index=True)

# Ensure is_augmented column exists
if 'is_augmented' not in df_balanced.columns:
    df_balanced['is_augmented'] = False

print(f"\n📊 Final Balanced Dataset Statistics:")
print(f"Total samples: {len(df_balanced)}")
print("\nFruit distribution:")
print(df_balanced['fruit'].value_counts())
print("\nFreshness distribution:")
print(df_balanced['fresh'].value_counts())

```

```

    if augmentation_stats:
        print(f"\n📊 Augmentation Statistics:")
        for fruit, stats in augmentation_stats.items():
            print(f"    {fruit}: {stats['original']} → {stats['total']} (+{stats['augmented']}) augmented)")

    return df_balanced

def create_stratified_splits(df):
    """
    Create 70/15/15 stratified splits - FIXED VERSION
    """
    print("\n📊 Creating Stratified Train/Val/Test Splits (70/15/15)...")

    if len(df) == 0:
        print("⚠️ Empty dataframe for splitting")
        return pd.DataFrame(), pd.DataFrame(), pd.DataFrame()

    # Create stratification key
    df['stratify_key'] = df['fruit'] + '_' + df['fresh'].astype(str)

    # Check if we have enough samples for stratification
    stratify_counts = df['stratify_key'].value_counts()
    min_count = stratify_counts.min()

    if min_count < 3: # Need at least 3 samples for 70/15/15 split
        print(f"⚠️ Some classes have too few samples ({min_count}). Using random splits...")
        train_df, temp_df = train_test_split(df, test_size=0.3, random_state=42)
        val_df, test_df = train_test_split(temp_df, test_size=0.5, random_state=42)
    else:
        try:
            # First split: 70% train, 30% temp
            train_df, temp_df = train_test_split(
                df, test_size=0.3, stratify=df['stratify_key'], random_state=42
            )

            # Second split: 15% val, 15% test from the 30% temp
            val_df, test_df = train_test_split(
                temp_df, test_size=0.5, stratify=temp_df['stratify_key'], random_state=42
            )

        except ValueError as e:
            print(f"⚠️ Stratification failed: {e}. Using random splits...")
            train_df, temp_df = train_test_split(df, test_size=0.3, random_state=42)
            val_df, test_df = train_test_split(temp_df, test_size=0.5, random_state=42)

    # Remove stratify_key column
    for split_df in [train_df, val_df, test_df]:
        if 'stratify_key' in split_df.columns:
            split_df.drop('stratify_key', axis=1, inplace=True)

```

```

print(f"✅ Train: {len(train_df)} ({len(train_df)/len(df)*100:.1f}%)")
print(f"✅ Val: {len(val_df)} ({len(val_df)/len(df)*100:.1f}%)")
print(f"✅ Test: {len(test_df)} ({len(test_df)/len(df)*100:.1f}%)")

return train_df, val_df, test_df

def compute_class_weights(train_df):
    """
    Compute class weights for both fruit and freshness tasks - FIXED VERSION
    """
    print("\n⚖️ Computing Class Weights...")

    if len(train_df) == 0:
        print("⚠️ Empty training dataframe, using default weights")
        return {}, {0: 1.0, 1: 1.0}

    try:
        # Fruit class weights
        fruit_classes = np.unique(train_df['fruit'])
        if len(fruit_classes) > 1:
            fruit_weights = compute_class_weight(
                'balanced',
                classes=fruit_classes,
                y=train_df['fruit']
            )
            fruit_class_weights = dict(zip(range(len(fruit_classes)), fruit_weights))
        else:
            fruit_class_weights = {0: 1.0}

        # Freshness class weights
        fresh_classes = np.unique(train_df['fresh'])
        if len(fresh_classes) > 1:
            fresh_weights = compute_class_weight(
                'balanced',
                classes=fresh_classes,
                y=train_df['fresh']
            )
            fresh_class_weights = dict(zip(fresh_classes, fresh_weights))
        else:
            fresh_class_weights = {0: 1.0, 1: 1.0}

        print(f"🍎 Fruit class weights: {fruit_class_weights}")
        print(f"🌱 Fresh class weights: {fresh_class_weights}")

        return fruit_class_weights, fresh_class_weights

    except Exception as e:
        print(f"⚠️ Error computing class weights: {e}")
        return {}, {0: 1.0, 1: 1.0}

#=====
# ENHANCED DATA GENERATORS - FIXED VERSION
#=====

class EnhancedMultiTaskDataGenerator(keras.utils.Sequence):

```

"""Enhanced data generator with augmentation support - FIXED VERSION"""

```

def __init__(self, dataframe, datagen, target_size=(224, 224), batch_size=
32,
            fruit_classes=None, shuffle=True):
    self.dataframe = dataframe.reset_index(drop=True)
    self.datagen = datagen
    self.target_size = target_size
    self.batch_size = batch_size
    self.fruit_classes = fruit_classes if fruit_classes is not None else
[]

    self.num_classes = len(self.fruit_classes)
    self.shuffle = shuffle
    self.indices = np.arange(len(self.dataframe))

    # Ensure we have valid data
    if len(self.dataframe) == 0:
        print("⚠ Warning: Empty dataframe provided to data generator")

    self.on_epoch_end()

def __len__(self):
    if len(self.dataframe) == 0:
        return 1 # Return at least 1 to avoid division by zero
    return int(np.ceil(len(self.dataframe) / self.batch_size))

def __getitem__(self, index):
    # Handle empty dataframe
    if len(self.dataframe) == 0:
        return self._get_dummy_batch()

    # Get batch indices
    start_idx = index * self.batch_size
    end_idx = min((index + 1) * self.batch_size, len(self.dataframe))
    batch_indices = self.indices[start_idx:end_idx]

    # Handle case where batch_indices is empty
    if len(batch_indices) == 0:
        return self._get_dummy_batch()

    batch_df = self.dataframe.iloc[batch_indices]

    batch_images = []
    batch_fruit_labels = []
    batch_fresh_labels = []

    for _, row in batch_df.iterrows():
        try:
            # Handle both real and augmented images
            is_augmented = row.get('is_augmented', False)

            if is_augmented and '_aug_' in str(row['filepath']):
                # Extract original filepath for augmented images
                original_path = str(row['filepath']).split('_aug_')[0]
                img_array = self._load_and_augment_image(original_path)
            else:
                img_array = self._load_image(row['filepath'])

```

```

        if img_array is not None:
            batch_images.append(img_array)

            # Get fruit label index
            if row['fruit'] in self.fruit_classes:
                fruit_idx = self.fruit_classes.index(row['fruit'])
            else:
                fruit_idx = 0 # Default to first class if not found

            batch_fruit_labels.append(fruit_idx)
            batch_fresh_labels.append(int(row['fresh']))

    except Exception as e:
        print(f"⚠️ Error processing row: {e}")
        # Add dummy data for failed samples
        dummy_img = self._get_dummy_image()
        batch_images.append(dummy_img)
        batch_fruit_labels.append(0)
        batch_fresh_labels.append(0)

    # Ensure we have at least one sample
    if len(batch_images) == 0:
        return self._get_dummy_batch()

    # Convert to arrays
    batch_images = np.array(batch_images)
    batch_fruit_labels = np.array(batch_fruit_labels)
    batch_fresh_labels = np.array(batch_fresh_labels)

    # Convert fruit labels to categorical
    if self.num_classes > 0:
        batch_fruit_labels = tf.keras.utils.to_categorical(
            batch_fruit_labels, num_classes=self.num_classes
        )
    else:
        # Handle case with no fruit classes
        batch_fruit_labels = np.zeros((len(batch_images), 1))

    return batch_images, {
        'fruit_output': batch_fruit_labels,
        'fresh_output': batch_fresh_labels
    }

def _get_dummy_batch(self):
    """Generate a dummy batch when no data is available"""
    dummy_images = np.random.rand(1, *self.target_size, 3) * 255
    if self.num_classes > 0:
        dummy_fruit = np.zeros((1, self.num_classes))
        dummy_fruit[0, 0] = 1 # Set first class to 1
    else:
        dummy_fruit = np.array([[0.0]])
    dummy_fresh = np.array([0.0])

    return dummy_images, {
        'fruit_output': dummy_fruit,
        'fresh_output': dummy_fresh
    }

```

```

    }

def _get_dummy_image(self):
    """Generate a dummy image"""
    return np.random.rand(*self.target_size, 3) * 255

def _load_image(self, filepath):
    """Load and preprocess a single image - FIXED VERSION"""
    try:
        if os.path.exists(filepath):
            img = tf.keras.preprocessing.image.load_img(
                filepath, target_size=self.target_size
            )
            img_array = tf.keras.preprocessing.image.img_to_array(img)
            img_array = self.datagen.standardize(img_array)
            return img_array
        else:
            # Return dummy image if file doesn't exist
            return self._get_dummy_image()
    except Exception as e:
        print(f"⚠ Error loading {filepath}: {e}")
        return self._get_dummy_image()

def _load_and_augment_image(self, filepath):
    """Load image and apply augmentation - FIXED VERSION"""
    try:
        if os.path.exists(filepath):
            img = tf.keras.preprocessing.image.load_img(
                filepath, target_size=self.target_size
            )
            img_array = tf.keras.preprocessing.image.img_to_array(img)

            # Apply random augmentation
            img_array = img_array.reshape((1,) + img_array.shape)

            # Create a flow and get one augmented image
            aug_iter = self.datagen.flow(img_array, batch_size=1)
            try:
                aug_img = next(aug_iter)[0] # FIXED: Use next() instead of
f .next()

            except (StopIteration, IndexError):
                # Fallback to original image if augmentation fails
                aug_img = img_array[0]

            return self.datagen.standardize(aug_img)
        else:
            return self._get_dummy_image()
    except Exception as e:
        print(f"⚠ Error augmenting {filepath}: {e}")
        return self._get_dummy_image()

def on_epoch_end(self):
    if self.shuffle:
        np.random.shuffle(self.indices)

#=====
# ENHANCED MODEL DEFINITIONS - FIXED VERSION

```

```

#=====

def create_multi_task_custom_cnn(input_shape=(224, 224, 3), num_fruit_classes=
9):
    """Custom CNN for Multi-Task Learning - FIXED VERSION"""

    # Ensure minimum number of classes
    num_fruit_classes = max(1, num_fruit_classes)

    inputs = layers.Input(shape=input_shape, name='input_layer')

    # Shared backbone with proper error handling
    x = layers.Conv2D(32, (3, 3), activation='relu', padding='same', name='con
v1_1')(inputs)
    x = layers.BatchNormalization(name='bn1_1')(x)
    x = layers.Conv2D(32, (3, 3), activation='relu', padding='same', name='con
v1_2')(x)
    x = layers.MaxPooling2D((2, 2), name='pool1')(x)
    x = layers.Dropout(0.25, name='dropout1')(x)

    x = layers.Conv2D(64, (3, 3), activation='relu', padding='same', name='con
v2_1')(x)
    x = layers.BatchNormalization(name='bn2_1')(x)
    x = layers.Conv2D(64, (3, 3), activation='relu', padding='same', name='con
v2_2')(x)
    x = layers.MaxPooling2D((2, 2), name='pool2')(x)
    x = layers.Dropout(0.25, name='dropout2')(x)

    x = layers.Conv2D(128, (3, 3), activation='relu', padding='same', name='co
nv3_1')(x)
    x = layers.BatchNormalization(name='bn3_1')(x)
    x = layers.Conv2D(128, (3, 3), activation='relu', padding='same', name='co
nv3_2')(x)
    x = layers.MaxPooling2D((2, 2), name='pool3')(x)
    x = layers.Dropout(0.25, name='dropout3')(x)

    x = layers.Conv2D(256, (3, 3), activation='relu', padding='same', name='co
nv4_1')(x)
    x = layers.BatchNormalization(name='bn4_1')(x)
    x = layers.Conv2D(256, (3, 3), activation='relu', padding='same', name='co
nv4_2')(x)
    x = layers.MaxPooling2D((2, 2), name='pool4')(x)
    x = layers.Dropout(0.25, name='dropout4')(x)

    # Shared dense Layers
    x = layers.Flatten(name='flatten')(x)
    x = layers.Dense(512, activation='relu', name='dense1')(x)
    x = layers.BatchNormalization(name='bn_dense1')(x)
    x = layers.Dropout(0.5, name='dropout_dense1')(x)
    shared_features = layers.Dense(256, activation='relu', name='shared_featur
es')(x)

    # Fruit classification head
    fruit_branch = layers.Dense(128, activation='relu', name='fruit_dense')(sh
ared_features)
    fruit_branch = layers.Dropout(0.3, name='fruit_dropout')(fruit_branch)
    fruit_output = layers.Dense(num_fruit_classes, activation='softmax', name

```

```

='fruit_output')(fruit_branch)

    # Freshness classification head
    fresh_branch = layers.Dense(64, activation='relu', name='fresh_dense')(shared_features)
    fresh_branch = layers.Dropout(0.3, name='fresh_dropout')(fresh_branch)
    fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')(fresh_branch)

    model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output],
name='Custom_CNN_MultiTask')
    return model

def create_multi_task_vgg16(input_shape=(224, 224, 3), num_fruit_classes=9):
    """VGG16 for Multi-Task Learning - FIXED VERSION"""
    num_fruit_classes = max(1, num_fruit_classes)

    try:
        base_model = VGG16(weights='imagenet', include_top=False, input_shape=input_shape)

        # Freeze early layers
        for layer in base_model.layers[:-4]:
            layer.trainable = False

        inputs = base_model.input
        x = base_model.output
        x = layers.GlobalAveragePooling2D(name='global_avg_pool')(x)
        x = layers.Dense(512, activation='relu', name='dense_512')(x)
        x = layers.BatchNormalization(name='bn_512')(x)
        x = layers.Dropout(0.5, name='dropout_512')(x)
        shared_features = layers.Dense(256, activation='relu', name='shared_256')(x)

        # Task-specific heads
        fruit_branch = layers.Dense(128, activation='relu', name='fruit_128')(shared_features)
        fruit_branch = layers.Dropout(0.3, name='fruit_dropout_vgg')(fruit_branch)
        fruit_output = layers.Dense(num_fruit_classes, activation='softmax', name='fruit_output')(fruit_branch)

        fresh_branch = layers.Dense(64, activation='relu', name='fresh_64')(shared_features)
        fresh_branch = layers.Dropout(0.3, name='fresh_dropout_vgg')(fresh_branch)
        fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')(fresh_branch)

        model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output], name='VGG16_MultiTask')
        return model
    except Exception as e:
        print(f"Error creating VGG16 model: {e}")
        return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

def create_multi_task_inceptionv3(input_shape=(224, 224, 3), num_fruit_classes

```

```

=9):
    """InceptionV3 for Multi-Task Learning - FIXED VERSION"""
    num_fruit_classes = max(1, num_fruit_classes)

    try:
        base_model = InceptionV3(weights='imagenet', include_top=False, input_shape=input_shape)

        for layer in base_model.layers[:-20]:
            layer.trainable = False

        inputs = base_model.input
        x = base_model.output
        x = layers.GlobalAveragePooling2D(name='global_avg_pool_inc')(x)
        x = layers.Dense(512, activation='relu', name='dense_512_inc')(x)
        x = layers.BatchNormalization(name='bn_512_inc')(x)
        x = layers.Dropout(0.5, name='dropout_512_inc')(x)
        shared_features = layers.Dense(256, activation='relu', name='shared_256_inc')(x)

        fruit_branch = layers.Dense(128, activation='relu', name='fruit_128_inc')(shared_features)
        fruit_branch = layers.Dropout(0.3, name='fruit_dropout_inc')(fruit_branch)
        fruit_output = layers.Dense(num_fruit_classes, activation='softmax', name='fruit_output')(fruit_branch)

        fresh_branch = layers.Dense(64, activation='relu', name='fresh_64_inc')(shared_features)
        fresh_branch = layers.Dropout(0.3, name='fresh_dropout_inc')(fresh_branch)
        fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')(fresh_branch)

        model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output], name='InceptionV3_MultiTask')
        return model
    except Exception as e:
        print(f"Error creating InceptionV3 model: {e}")
        return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

def create_multi_task_resnet152(input_shape=(224, 224, 3), num_fruit_classes=9):
    """ResNet152V2 for Multi-Task Learning - FIXED VERSION"""
    num_fruit_classes = max(1, num_fruit_classes)

    try:
        base_model = ResNet152V2(weights='imagenet', include_top=False, input_shape=input_shape)

        for layer in base_model.layers[:-30]:
            layer.trainable = False

        inputs = base_model.input
        x = base_model.output
        x = layers.GlobalAveragePooling2D(name='global_avg_pool_res')(x)
        x = layers.Dense(512, activation='relu', name='dense_512_res')(x)

```

```

x = layers.BatchNormalization(name='bn_512_res')(x)
x = layers.Dropout(0.5, name='dropout_512_res')(x)
shared_features = layers.Dense(256, activation='relu', name='shared_25
6_res')(x)

    fruit_branch = layers.Dense(128, activation='relu', name='fruit_128_re
s')(shared_features)
    fruit_branch = layers.Dropout(0.3, name='fruit_dropout_res')(fruit_bra
nch)
    fruit_output = layers.Dense(num_fruit_classes, activation='softmax', n
ame='fruit_output')(fruit_branch)

    fresh_branch = layers.Dense(64, activation='relu', name='fresh_64_re
s')(shared_features)
    fresh_branch = layers.Dropout(0.3, name='fresh_dropout_res')(fresh_bra
nch)
    fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_outpu
t')(fresh_branch)

    model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_outpu
t], name='ResNet152_MultiTask')
    return model
except Exception as e:
    print(f"Error creating ResNet152 model: {e}")
    return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

def create_darknet_conv_block(x, filters, kernel_size=3, strides=1):
    """Darknet convolutional block - FIXED VERSION"""
    x = layers.Conv2D(filters, kernel_size, strides=strides, padding='same', u
se_bias=False)(x)
    x = layers.BatchNormalization()(x)
    x = layers.LeakyReLU(alpha=0.1)(x)
    return x

def create_darknet_residual_block(x, filters):
    """Darknet residual block - FIXED VERSION"""
    shortcut = x
    x = create_darknet_conv_block(x, filters//2, 1)
    x = create_darknet_conv_block(x, filters, 3)
    x = layers.Add()([shortcut, x])
    return x

def create_multi_task_yolov3_backbone(input_shape=(224, 224, 3), num_fruit_cla
sses=9):
    """YOLOv3 Darknet-53 backbone for Multi-Task Learning - FIXED VERSION"""
    num_fruit_classes = max(1, num_fruit_classes)

    inputs = layers.Input(shape=input_shape, name='darknet_input')

    # Darknet-53 backbone (simplified)
    x = create_darknet_conv_block(inputs, 32, 3)
    x = create_darknet_conv_block(x, 64, 3, strides=2)
    x = create_darknet_residual_block(x, 64)

    x = create_darknet_conv_block(x, 128, 3, strides=2)
    for _ in range(2):
        x = create_darknet_residual_block(x, 128)

```

```

x = create_darknet_conv_block(x, 256, 3, strides=2)
for _ in range(4):
    x = create_darknet_residual_block(x, 256)

x = create_darknet_conv_block(x, 512, 3, strides=2)
for _ in range(4):
    x = create_darknet_residual_block(x, 512)

# Shared features
x = layers.GlobalAveragePooling2D(name='darknet_global_pool')(x)
x = layers.Dense(128, activation='relu', kernel_regularizer=regularizers.l
2(0.01), name='darknet_dense_128')(x)
x = layers.Dropout(0.5, name='darknet_dropout_128')(x)
shared_features = layers.Dense(64, activation='relu', kernel_regularizer=r
egularizers.l2(0.01), name='darknet_shared_64')(x)

# Task-specific heads
fruit_output = layers.Dense(num_fruit_classes, activation='softmax', name
='fruit_output')(shared_features)
fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')
(shared_features)

model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output],
name='YOLOv3_MultiTask')
return model

def create_multi_task_efficientnet(input_shape=(224, 224, 3), num_fruit_classe
s=9):
    """EfficientNetB4 for Multi-Task Learning"""
    from tensorflow.keras.applications import EfficientNetB4

    num_fruit_classes = max(1, num_fruit_classes)

    try:
        base_model = EfficientNetB4(weights='imagenet', include_top=False, inp
ut_shape=input_shape)

        # Fine-tune the top layers
        for layer in base_model.layers[:-20]:
            layer.trainable = False

        inputs = base_model.input
        x = base_model.output
        x = layers.GlobalAveragePooling2D(name='global_avg_pool_eff')(x)
        x = layers.Dense(512, activation='relu', name='dense_512_eff')(x)
        x = layers.BatchNormalization(name='bn_512_eff')(x)
        x = layers.Dropout(0.3, name='dropout_512_eff')(x)
        shared_features = layers.Dense(256, activation='relu', name='shared_25
6_eff')(x)

        # Task-specific heads
        fruit_branch = layers.Dense(128, activation='relu', name='fruit_128_ef
f')(shared_features)
        fruit_branch = layers.Dropout(0.2, name='fruit_dropout_eff')(fruit_bra
nch)
        fruit_output = layers.Dense(num_fruit_classes, activation='softmax', n

```

```

ame='fruit_output')(fruit_branch)

        fresh_branch = layers.Dense(64, activation='relu', name='fresh_64_eff')(shared_features)
        fresh_branch = layers.Dropout(0.2, name='fresh_dropout_eff')(fresh_branch)
        fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')(fresh_branch)

        model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output], name='EfficientNetB4_MultiTask')
        return model
    except Exception as e:
        print(f"Error creating EfficientNet model: {e}")
        return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

def create_multi_task_densenet(input_shape=(224, 224, 3), num_fruit_classes=9):
    """DenseNet121 for Multi-Task Learning"""
    from tensorflow.keras.applications import DenseNet121

    num_fruit_classes = max(1, num_fruit_classes)

    try:
        base_model = DenseNet121(weights='imagenet', include_top=False, input_shape=input_shape)

        # Fine-tune Later Layers
        for layer in base_model.layers[:-15]:
            layer.trainable = False

        inputs = base_model.input
        x = base_model.output
        x = layers.GlobalAveragePooling2D(name='global_avg_pool_dense')(x)
        x = layers.Dense(512, activation='relu', name='dense_512_dense')(x)
        x = layers.BatchNormalization(name='bn_512_dense')(x)
        x = layers.Dropout(0.4, name='dropout_512_dense')(x)
        shared_features = layers.Dense(256, activation='relu', name='shared_256_dense')(x)

        # Task-specific heads
        fruit_branch = layers.Dense(128, activation='relu', name='fruit_128_dense')(shared_features)
        fruit_branch = layers.Dropout(0.3, name='fruit_dropout_dense')(fruit_branch)
        fruit_output = layers.Dense(num_fruit_classes, activation='softmax', name='fruit_output')(fruit_branch)

        fresh_branch = layers.Dense(64, activation='relu', name='fresh_64_dense')(shared_features)
        fresh_branch = layers.Dropout(0.3, name='fresh_dropout_dense')(fresh_branch)
        fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')(fresh_branch)

        model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output], name='DenseNet121_MultiTask')

```

```

        return model
    except Exception as e:
        print(f"Error creating DenseNet model: {e}")
        return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

def create_multi_task_convnext_base(input_shape=(224, 224, 3), num_fruit_classes=9):
    """ConvNeXt-Base Pretrained for Multi-Task Learning"""
    from tensorflow.keras.applications import ConvNeXtBase

    num_fruit_classes = max(1, num_fruit_classes)

    try:
        base_model = ConvNeXtBase(weights='imagenet', include_top=False, input_shape=input_shape)

        for layer in base_model.layers[:-10]:
            layer.trainable = False

        inputs = base_model.input
        x = base_model.output
        x = layers.GlobalAveragePooling2D(name='convnext_gap')(x)
        x = layers.Dense(512, activation='relu', name='convnext_dense_512')(x)
        x = layers.BatchNormalization(name='convnext_bn_512')(x)
        x = layers.Dropout(0.3, name='convnext_dropout_512')(x)
        shared_features = layers.Dense(256, activation='relu', name='convnext_shared_256')(x)

        # Task-specific heads
        fruit_branch = layers.Dense(128, activation='relu', name='convnext_fruit_128')(shared_features)
        fruit_branch = layers.Dropout(0.2, name='convnext_fruit_dropout')(fruit_branch)
        fruit_output = layers.Dense(num_fruit_classes, activation='softmax', name='fruit_output')(fruit_branch)

        fresh_branch = layers.Dense(64, activation='relu', name='convnext_fresh_64')(shared_features)
        fresh_branch = layers.Dropout(0.2, name='convnext_fresh_dropout')(fresh_branch)
        fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_output')(fresh_branch)

        model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_output], name='ConvNeXtBase_MultiTask')
        return model
    except Exception as e:
        print(f"Error creating ConvNeXt model: {e}")
        return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

def create_multi_task_vit_base_pretrained(input_shape=(224, 224, 3), num_fruit_classes=9):
    """Google's Pretrained ViT-Base-16 for Multi-Task Learning - Best ViT Choice"""

    num_fruit_classes = max(1, num_fruit_classes)

```

```

try:
    # Load pretrained ViT-Base-16 from TensorFlow Hub
    vit_base_url = "https://tfhub.dev/google/vit_base_patch16_224/1"
    base_model = hub.KerasLayer(vit_base_url, trainable=True)

    inputs = layers.Input(shape=input_shape, name='vit_input')

    # Pretrained ViT feature extraction (this has ImageNet-21k pretrainin
g!)
    x = base_model(inputs) # Shape: (batch_size, 768)

    # Add custom multi-task heads optimized for your fruit/freshness tasks
    x = layers.Dense(512, activation='relu', name='vit_dense_512')(x)
    x = layers.BatchNormalization(name='vit_bn_512')(x)
    x = layers.Dropout(0.3, name='vit_dropout_512')(x)
    shared_features = layers.Dense(256, activation='relu', name='vit_share
d_256')(x)

    # Task-specific heads
    fruit_branch = layers.Dense(128, activation='relu', name='vit_fruit_12
8')(shared_features)
    fruit_branch = layers.Dropout(0.2, name='vit_fruit_dropout')(fruit_bra
nch)
    fruit_output = layers.Dense(num_fruit_classes, activation='softmax', n
ame='fruit_output')(fruit_branch)

    fresh_branch = layers.Dense(64, activation='relu', name='vit_fresh_6
4')(shared_features)
    fresh_branch = layers.Dropout(0.2, name='vit_fresh_dropout')(fresh_bra
nch)
    fresh_output = layers.Dense(1, activation='sigmoid', name='fresh_outpu
t')(fresh_branch)

    model = models.Model(inputs=inputs, outputs=[fruit_output, fresh_outpu
t], name='ViT_Base_Pretrained_MultiTask')
    return model
except Exception as e:
    print(f"Error creating pretrained ViT model: {e}")
    return create_multi_task_custom_cnn(input_shape, num_fruit_classes)

#=====
# TRAINING AND EVALUATION FUNCTIONS - FIXED VERSION
#=====

def cleanup_memory():
    """Clean up memory between model training"""
    import gc
    try:
        gc.collect()
        tf.keras.backend.clear_session()
    except:
        pass

def compile_and_train_multi_task_model(model, model_name, train_gen, val_gen,
epochs=15, alpha=0.7, class_weights=None):
    """Compile and train a multi-task model with class weights - FIXED VERSIO
N"""

```

```

print(f"\n{'='*50}")
print(f"Training {model_name}")
print(f"{'='*50}")

# Check if generators have data
if len(train_gen) == 0 or len(val_gen) == 0:
    print("⚠ Warning: Empty data generators")
    # Create dummy history
    dummy_history = type('DummyHistory', (), {})(())
    dummy_history.history = {
        'loss': [1.0] * epochs,
        'val_loss': [1.0] * epochs,
        'fruit_output_accuracy': [0.1] * epochs,
        'val_fruit_output_accuracy': [0.1] * epochs,
        'fresh_output_accuracy': [0.5] * epochs,
        'val_fresh_output_accuracy': [0.5] * epochs,
        'fruit_output_loss': [1.0] * epochs,
        'val_fruit_output_loss': [1.0] * epochs,
        'fresh_output_loss': [0.5] * epochs,
        'val_fresh_output_loss': [0.5] * epochs
    }
    return model, dummy_history

# Enhanced Learning rate for different models
if 'ViT' in model_name or 'Vision' in model_name:
    learning_rate = 0.0005
elif 'ResNet' in model_name:
    learning_rate = 0.0005
else:
    learning_rate = 0.001

try:
    model.compile(
        optimizer=optimizers.Adam(learning_rate=learning_rate),
        loss={
            'fruit_output': 'categorical_crossentropy',
            'fresh_output': 'binary_crossentropy'
        },
        loss_weights={
            'fruit_output': alpha,
            'fresh_output': 1 - alpha
        },
        metrics={
            'fruit_output': ['accuracy'],
            'fresh_output': ['accuracy']
        }
    )

    callbacks = [
        EarlyStopping(
            monitor='val_loss',
            patience=5, # Reduced patience for faster training
            restore_best_weights=True,
            verbose=1
        ),
        ReduceLROnPlateau(
            monitor='val_loss',

```

```

        factor=0.5,
        patience=3,
        min_lr=1e-7,
        verbose=1
    )
]

# Train the model with proper error handling
history = model.fit(
    train_gen,
    epochs=epochs,
    validation_data=val_gen,
    callbacks=callbacks,
    verbose=1,
)

return model, history

except Exception as e:
    print(f"Error during training {model_name}: {e}")
    # Create dummy history
    dummy_history = type('DummyHistory', (), {})(())
    dummy_history.history = {
        'loss': [1.0] * epochs,
        'val_loss': [1.0] * epochs,
        'fruit_output_accuracy': [0.1] * epochs,
        'val_fruit_output_accuracy': [0.1] * epochs,
        'fresh_output_accuracy': [0.5] * epochs,
        'val_fresh_output_accuracy': [0.5] * epochs,
        'fruit_output_loss': [1.0] * epochs,
        'val_fruit_output_loss': [1.0] * epochs,
        'fresh_output_loss': [0.5] * epochs,
        'val_fresh_output_loss': [0.5] * epochs
    }
    return model, dummy_history

def evaluate_multi_task_model(model, train_gen, val_gen, test_gen, model_name):
    """Evaluate multi-task model on all three sets - FIXED VERSION"""
    print(f"Evaluating {model_name}...")

    def safe_evaluate(generator, split_name):
        """Safely evaluate model on a generator"""
        try:
            if len(generator) == 0:
                print(f"⚠ Empty {split_name} generator")
                return {'loss': 999.0, 'fruit_output_accuracy': 0.0, 'fresh_output_accuracy': 0.0}

            results = model.evaluate(generator, verbose=0, return_dict=True)
            return results
        except Exception as e:
            print(f"⚠ Error evaluating {split_name}: {e}")
            return {'loss': 999.0, 'fruit_output_accuracy': 0.0, 'fresh_output_accuracy': 0.0}

    train_results = safe_evaluate(train_gen, 'train')

```

```

val_results = safe_evaluate(val_gen, 'validation')
test_results = safe_evaluate(test_gen, 'test')

results = {
    'Model': model_name,
    'Train Loss': train_results.get('loss', 999.0),
    'Valid Loss': val_results.get('loss', 999.0),
    'Test Loss': test_results.get('loss', 999.0),
    'Train Fruit Acc': train_results.get('fruit_output_accuracy', 0.0),
    'Valid Fruit Acc': val_results.get('fruit_output_accuracy', 0.0),
    'Test Fruit Acc': test_results.get('fruit_output_accuracy', 0.0),
    'Train Fresh Acc': train_results.get('fresh_output_accuracy', 0.0),
    'Valid Fresh Acc': val_results.get('fresh_output_accuracy', 0.0),
    'Test Fresh Acc': test_results.get('fresh_output_accuracy', 0.0)
}

return results

#=====
# ALL VISUALIZATION FUNCTIONS - FIXED FOR 9 MODELS
#=====

def save_individual_fruit_accuracy_comparison(results_df, output_dir="conference_figures"):
    """Save individual fruit classification accuracy comparison"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    plt.figure(figsize=(12, 8))

    models = results_df['Model']
    x_pos = np.arange(len(models))

    # FIXED: Updated colors for 9 models
    colors = ['#FF6B6B', '#4ECDC4', '#45B7D1', '#96CEB4', '#FECA57', '#FF9FF3', '#A8E6CF', '#FFD93D', '#FF8A80']
    bars = plt.bar(x_pos, results_df['Test Fruit Acc'], color=colors[:len(models)], alpha=0.8, edgecolor='black', linewidth=1)

    for i, bar in enumerate(bars):
        height = bar.get_height()
        plt.text(bar.get_x() + bar.get_width()/2., height + 0.01,
                 f'{height:.3f}', ha='center', va='bottom', fontweight='bold',
                 fontsize=11)

    plt.title('Fruit Classification Test Accuracy Comparison', fontsize=16, fontweight='bold', pad=20)
    plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
    plt.ylabel('Test Accuracy', fontsize=14, fontweight='bold')
    plt.xticks(x_pos, models, rotation=45, ha='right', fontsize=12)
    plt.ylim(0, 1.1)
    plt.grid(True, alpha=0.3, axis='y')
    plt.tight_layout()

    filename = os.path.join(output_dir, 'fruit_classification_accuracy_comparison.png')
    plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')

```

```

plt.show()
print(f"✓ Fruit accuracy comparison saved: {filename}")

def save_individual_freshness_accuracy_comparison(results_df, output_dir="conference_figures"):
    """Save individual freshness classification accuracy comparison"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    plt.figure(figsize=(12, 8))

    models = results_df['Model']
    x_pos = np.arange(len(models))

    # FIXED: Updated colors for 9 models
    colors = ['#FF6B6B', '#4ECDC4', '#45B7D1', '#96CEB4', '#FECA57', '#FF9FF3', '#A8E6CF', '#FFD93D', '#FF8A80']
    bars = plt.bar(x_pos, results_df['Test Fresh Acc'], color=colors[:len(models)], alpha=0.8, edgecolor='black', linewidth=1)

    for i, bar in enumerate(bars):
        height = bar.get_height()
        plt.text(bar.get_x() + bar.get_width()/2., height + 0.01,
                 f'{height:.3f}', ha='center', va='bottom', fontweight='bold',
                 fontsize=11)

    plt.title('Freshness Classification Test Accuracy Comparison', fontsize=16, fontweight='bold', pad=20)
    plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
    plt.ylabel('Test Accuracy', fontsize=14, fontweight='bold')
    plt.xticks(x_pos, models, rotation=45, ha='right', fontsize=12)
    plt.ylim(0, 1.1)
    plt.grid(True, alpha=0.3, axis='y')
    plt.tight_layout()

    filename = os.path.join(output_dir, 'freshness_classification_accuracy_comparison.png')
    plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
    plt.show()
    print(f"✓ Freshness accuracy comparison saved: {filename}")

def save_individual_combined_performance_plot(results_df, output_dir="conference_figures"):
    """Save individual combined performance comparison"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    plt.figure(figsize=(12, 8))

    x = range(len(results_df))
    plt.plot(x, results_df['Test Fruit Acc'], 'o-', label='Fruit Classification',
             linewidth=3, markersize=10, color='#FF6B6B')
    plt.plot(x, results_df['Test Fresh Acc'], 's-', label='Freshness Classification',
             linewidth=3, markersize=10, color='#4ECDC4')

```

```

plt.title('Multi-Task Test Performance Comparison', fontsize=16, fontweight='bold', pad=20)
plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
plt.ylabel('Test Accuracy', fontsize=14, fontweight='bold')
plt.xticks(x, results_df['Model'], rotation=45, ha='right', fontsize=12)
plt.legend(fontsize=12, loc='best')
plt.grid(True, alpha=0.3)
plt.ylim(0, 1.1)
plt.tight_layout()

filename = os.path.join(output_dir, 'combined_performance_comparison.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Combined performance comparison saved: {filename}")

def save_individual_loss_comparison(results_df, output_dir="conference_figures"):
    """Save individual loss comparison"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    plt.figure(figsize=(12, 8))

    x = range(len(results_df))
    plt.plot(x, results_df['Train Loss'], 'o-', label='Training Loss',
             linewidth=3, markersize=10, color='#45B7D1')
    plt.plot(x, results_df['Valid Loss'], 's-', label='Validation Loss',
             linewidth=3, markersize=10, color='#96CEB4')
    plt.plot(x, results_df['Test Loss'], '^--', label='Test Loss',
             linewidth=3, markersize=10, color='#FECA57')

    plt.title('Multi-Task Loss Comparison Across Models', fontsize=16, fontweight='bold', pad=20)
    plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
    plt.ylabel('Loss Value', fontsize=14, fontweight='bold')
    plt.xticks(x, results_df['Model'], rotation=45, ha='right', fontsize=12)
    plt.legend(fontsize=12, loc='best')
    plt.grid(True, alpha=0.3)
    plt.tight_layout()

    filename = os.path.join(output_dir, 'loss_comparison.png')
    plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
    plt.show()
    print(f"✓ Loss comparison saved: {filename}")

def save_individual_performance_heatmaps(results_df, output_dir="conference_figures"):
    """Save individual performance heatmaps"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    # Fruit Classification Heatmap
    plt.figure(figsize=(14, 6))
    fruit_data = results_df[['Train Fruit Acc', 'Valid Fruit Acc', 'Test Fruit Acc']].T
    fruit_data.columns = results_df['Model']

```

```

sns.heatmap(fruit_data, annot=True, cmap='Oranges', fmt='.3f',
             cbar_kws={'label': 'Accuracy'}, square=False, linewidths=0.5)
plt.title('Fruit Classification Performance Heatmap', fontsize=16, fontwei
ght='bold', pad=20)
plt.ylabel('Dataset Split', fontsize=14, fontweight='bold')
plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()

filename = os.path.join(output_dir, 'fruit_classification_heatmap.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Fruit classification heatmap saved: {filename}")

# Freshness Classification Heatmap
plt.figure(figsize=(14, 6))
fresh_data = results_df[['Train Fresh Acc', 'Valid Fresh Acc', 'Test Fresh
Acc']].T
fresh_data.columns = results_df['Model']

sns.heatmap(fresh_data, annot=True, cmap='Blues', fmt='.3f',
             cbar_kws={'label': 'Accuracy'}, square=False, linewidths=0.5)
plt.title('Freshness Classification Performance Heatmap', fontsize=16, fon
tweight='bold', pad=20)
plt.ylabel('Dataset Split', fontsize=14, fontweight='bold')
plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()

filename = os.path.join(output_dir, 'freshness_classification_heatmap.pn
g')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Freshness classification heatmap saved: {filename}")

def save_individual_correlation_scatter(results_df, output_dir="conference_fig
ures"):
    """Save individual correlation scatter plot"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    plt.figure(figsize=(10, 8))

    # FIXED: Updated colors for 9 models
    colors = ['#FF6B6B', '#4ECDC4', '#45B7D1', '#96CEB4', '#FECA57', '#FF9FF
3', '#A8E6CF', '#FFD93D', '#FF8A80']
    scatter = plt.scatter(results_df['Test Fruit Acc'], results_df['Test Fresh
Acc'],
                          s=200, c=colors[:len(results_df)], alpha=0.8, edgecol
ors='black', linewidth=2)

    # Add model labels
    for i, model in enumerate(results_df['Model']):
        plt.annotate(model, (results_df['Test Fruit Acc'].iloc[i], results_df
['Test Fresh Acc'].iloc[i]),
                      xytext=(10, 10), textcoords='offset points', fontsize=11,
                      fontweight='bold',

```

```

        bbox=dict(boxstyle='round,pad=0.3', facecolor='white', alp
ha=0.8))

    plt.xlabel('Fruit Classification Test Accuracy', fontsize=14, fontweight
='bold')
    plt.ylabel('Freshness Classification Test Accuracy', fontsize=14, fontweig
ht='bold')
    plt.title('Multi-Task Performance Correlation Analysis', fontsize=16, font
weight='bold', pad=20)
    plt.grid(True, alpha=0.3)
    plt.xlim(0, 1.05)
    plt.ylim(0, 1.05)
    plt.tight_layout()

    filename = os.path.join(output_dir, 'performance_correlation_scatter.png')
    plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
    plt.show()
    print(f"✓ Performance correlation scatter saved: {filename}")

def save_individual_average_performance_bar(results_df, output_dir="conference
_figures"):
    """Save individual average performance bar chart"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    plt.figure(figsize=(12, 8))

    avg_fruit = (results_df['Train Fruit Acc'] + results_df['Valid Fruit Acc']
+ results_df['Test Fruit Acc']) / 3
    avg_fresh = (results_df['Train Fresh Acc'] + results_df['Valid Fresh Acc']
+ results_df['Test Fresh Acc']) / 3

    x = np.arange(len(results_df))
    width = 0.35

    bars1 = plt.bar(x - width/2, avg_fruit, width, label='Fruit Classificatio
n',
                    alpha=0.8, color='#FF6B6B', edgecolor='black', linewidth=1)
    bars2 = plt.bar(x + width/2, avg_fresh, width, label='Freshness Classifica
tion',
                    alpha=0.8, color='#4ECDC4', edgecolor='black', linewidth=1)

    # Add value labels
    for bars in [bars1, bars2]:
        for bar in bars:
            height = bar.get_height()
            plt.text(bar.get_x() + bar.get_width()/2., height + 0.01,
                     f'{height:.3f}', ha='center', va='bottom', fontweight='bol
d', fontsize=10)

    plt.title('Average Performance Across All Dataset Splits', fontsize=16, fo
ntweight='bold', pad=20)
    plt.xlabel('Deep Learning Models', fontsize=14, fontweight='bold')
    plt.ylabel('Average Accuracy', fontsize=14, fontweight='bold')
    plt.xticks(x, results_df['Model'], rotation=45, ha='right', fontsize=12)
    plt.legend(fontsize=12, loc='best')
    plt.grid(True, alpha=0.3, axis='y')

```

```

plt.ylim(0, 1.1)
plt.tight_layout()

filename = os.path.join(output_dir, 'average_performance_comparison.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Average performance comparison saved: {filename}")

def save_individual_training_history(history, model_name, output_dir="conference_figures"):
    """Save individual training history plots"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    if history is None or not hasattr(history, 'history'):
        print(f"No training history available for {model_name}")
        return

    # Create 2x2 subplot for training history
    fig, axes = plt.subplots(2, 2, figsize=(15, 12))
    fig.suptitle(f'{model_name} - Training History Analysis', fontsize=18, fontweight='bold')

    # 1. Fruit Classification Accuracy
    if 'fruit_output_accuracy' in history.history:
        axes[0,0].plot(history.history['fruit_output_accuracy'], 'b-', label='Training', linewidth=2)
        axes[0,0].plot(history.history['val_fruit_output_accuracy'], 'r-', label='Validation', linewidth=2)
        axes[0,0].set_title('Fruit Classification Accuracy', fontsize=14, fontweight='bold')
        axes[0,0].set_xlabel('Epochs', fontsize=12)
        axes[0,0].set_ylabel('Accuracy', fontsize=12)
        axes[0,0].legend(fontsize=11)
        axes[0,0].grid(True, alpha=0.3)

    # 2. Freshness Classification Accuracy
    if 'fresh_output_accuracy' in history.history:
        axes[0,1].plot(history.history['fresh_output_accuracy'], 'b-', label='Training', linewidth=2)
        axes[0,1].plot(history.history['val_fresh_output_accuracy'], 'r-', label='Validation', linewidth=2)
        axes[0,1].set_title('Freshness Classification Accuracy', fontsize=14, fontweight='bold')
        axes[0,1].set_xlabel('Epochs', fontsize=12)
        axes[0,1].set_ylabel('Accuracy', fontsize=12)
        axes[0,1].legend(fontsize=11)
        axes[0,1].grid(True, alpha=0.3)

    # 3. Fruit Classification Loss
    if 'fruit_output_loss' in history.history:
        axes[1,0].plot(history.history['fruit_output_loss'], 'b-', label='Training', linewidth=2)
        axes[1,0].plot(history.history['val_fruit_output_loss'], 'r-', label='Validation', linewidth=2)
        axes[1,0].set_title('Fruit Classification Loss', fontsize=14, fontweight='bold')

```

```

axes[1,0].set_xlabel('Epochs', fontsize=12)
axes[1,0].set_ylabel('Loss', fontsize=12)
axes[1,0].legend(fontsize=11)
axes[1,0].grid(True, alpha=0.3)

# 4. Freshness Classification Loss
if 'fresh_output_loss' in history.history:
    axes[1,1].plot(history.history['fresh_output_loss'], 'b-', label='Training', linewidth=2)
    axes[1,1].plot(history.history['val_fresh_output_loss'], 'r-', label='Validation', linewidth=2)
    axes[1,1].set_title('Freshness Classification Loss', fontsize=14, fontweight='bold')
    axes[1,1].set_xlabel('Epochs', fontsize=12)
    axes[1,1].set_ylabel('Loss', fontsize=12)
    axes[1,1].legend(fontsize=11)
    axes[1,1].grid(True, alpha=0.3)

plt.tight_layout()

safe_model_name = model_name.lower().replace(' ', '_').replace('+', '_').replace('-', '_')
filename = os.path.join(output_dir, f'{safe_model_name}_training_history.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Training history saved: {filename}")

def save_best_model_training_history(history, model_name, output_dir="conference_figures"):
    """Save best model training history plots"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    if history is None or not hasattr(history, 'history'):
        print(f"No training history available for best model")
        return

    # Create 2x2 subplot for training history
    fig, axes = plt.subplots(2, 2, figsize=(15, 12))
    fig.suptitle(f'Best Model ({model_name}) - Training History Analysis', fontsize=18, fontweight='bold')

    # 1. Fruit Classification Accuracy
    if 'fruit_output_accuracy' in history.history:
        axes[0,0].plot(history.history['fruit_output_accuracy'], 'b-', label='Training', linewidth=2)
        axes[0,0].plot(history.history['val_fruit_output_accuracy'], 'r-', label='Validation', linewidth=2)
        axes[0,0].set_title('Fruit Classification Accuracy', fontsize=14, fontweight='bold')
        axes[0,0].set_xlabel('Epochs', fontsize=12)
        axes[0,0].set_ylabel('Accuracy', fontsize=12)
        axes[0,0].legend(fontsize=11)
        axes[0,0].grid(True, alpha=0.3)

    # 2. Freshness Classification Accuracy

```

```

    if 'fresh_output_accuracy' in history.history:
        axes[0,1].plot(history.history['fresh_output_accuracy'], 'b-', label
='Training', linewidth=2)
        axes[0,1].plot(history.history['val_fresh_output_accuracy'], 'r-', lab
el='Validation', linewidth=2)
        axes[0,1].set_title('Freshness Classification Accuracy', fontsize=14,
fontweight='bold')
        axes[0,1].set_xlabel('Epochs', fontsize=12)
        axes[0,1].set_ylabel('Accuracy', fontsize=12)
        axes[0,1].legend(fontsize=11)
        axes[0,1].grid(True, alpha=0.3)

    # 3. Fruit Classification Loss
    if 'fruit_output_loss' in history.history:
        axes[1,0].plot(history.history['fruit_output_loss'], 'b-', label='Trai
ning', linewidth=2)
        axes[1,0].plot(history.history['val_fruit_output_loss'], 'r-', label
='Validation', linewidth=2)
        axes[1,0].set_title('Fruit Classification Loss', fontsize=14, fontweig
ht='bold')
        axes[1,0].set_xlabel('Epochs', fontsize=12)
        axes[1,0].set_ylabel('Loss', fontsize=12)
        axes[1,0].legend(fontsize=11)
        axes[1,0].grid(True, alpha=0.3)

    # 4. Freshness Classification Loss
    if 'fresh_output_loss' in history.history:
        axes[1,1].plot(history.history['fresh_output_loss'], 'b-', label='Trai
ning', linewidth=2)
        axes[1,1].plot(history.history['val_fresh_output_loss'], 'r-', label
='Validation', linewidth=2)
        axes[1,1].set_title('Freshness Classification Loss', fontsize=14, font
weight='bold')
        axes[1,1].set_xlabel('Epochs', fontsize=12)
        axes[1,1].set_ylabel('Loss', fontsize=12)
        axes[1,1].legend(fontsize=11)
        axes[1,1].grid(True, alpha=0.3)

plt.tight_layout()

filename = os.path.join(output_dir, 'best_training_history.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Best model training history saved: {filename}")

def save_best_model_confusion_matrices(model, test_gen, model_name, fruit_clas
ses, output_dir="conference_figures"):
    """Save confusion matrices for best model - both tasks"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    print(f"Creating confusion matrices for best model ({model_name})...")

    try:
        # Get predictions
        all_predictions = []
        all_fruit_labels = []

```

```

all_fresh_labels = []

for i in range(len(test_gen)):
    batch_x, batch_y = test_gen[i]
    predictions = model.predict(batch_x, verbose=0)
    fruit_preds = predictions[0]
    fresh_preds = predictions[1]

    all_predictions.append((fruit_preds, fresh_preds))
    all_fruit_labels.extend(np.argmax(batch_y['fruit_output'], axis=
1))

    all_fresh_labels.extend(batch_y['fresh_output'])

# Combine predictions
fruit_predictions = np.vstack([pred[0] for pred in all_predictions])
fresh_predictions = np.vstack([pred[1] for pred in all_predictions])
fruit_pred_classes = np.argmax(fruit_predictions, axis=1)
fresh_pred_classes = (fresh_predictions > 0.5).astype(int).flatten()

# 1. Best Model Fruit Classification Confusion Matrix
plt.figure(figsize=(12, 10))
fruit_cm = confusion_matrix(all_fruit_labels, fruit_pred_classes)

display_fruit_classes = [fruit.capitalize() for fruit in fruit_classes]

sns.heatmap(fruit_cm, annot=True, fmt='d', cmap='Oranges',
            xticklabels=display_fruit_classes,
            yticklabels=display_fruit_classes,
            square=True, linewidths=0.5)
plt.title(f'Best Model ({model_name}) - Fruit Classification Confusion
Matrix',
        fontsize=16, fontweight='bold', pad=20)
plt.xlabel('Predicted Fruit Class', fontsize=14, fontweight='bold')
plt.ylabel('Actual Fruit Class', fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()

filename = os.path.join(output_dir, 'best_fruit_confusion_matrix.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Best model fruit confusion matrix saved: {filename}")

# 2. Best Model Freshness Classification Confusion Matrix
plt.figure(figsize=(8, 6))
fresh_cm = confusion_matrix(all_fresh_labels, fresh_pred_classes)

sns.heatmap(fresh_cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Fresh', 'Rotten'],
            yticklabels=['Fresh', 'Rotten'],
            square=True, linewidths=0.5)
plt.title(f'Best Model ({model_name}) - Freshness Classification Confu
sion Matrix',
        fontsize=16, fontweight='bold', pad=20)
plt.xlabel('Predicted Freshness Class', fontsize=14, fontweight='bol
d')

```

```

plt.ylabel('Actual Freshness Class', fontsize=14, fontweight='bold')
plt.tight_layout()

filename = os.path.join(output_dir, 'best_freshness_confusion_matrix.png')

plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Best model freshness confusion matrix saved: {filename}")

except Exception as e:
    print(f"Error creating confusion matrices for best model: {e}")

def save_best_model_roc_curve(model, test_gen, model_name, output_dir="conference_figures"):
    """Save ROC curve for best model freshness classification"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    print(f"Creating ROC curve for best model ({model_name})...")

    try:
        # Get predictions for freshness task
        all_fresh_labels = []
        all_fresh_probs = []

        for i in range(len(test_gen)):
            batch_x, batch_y = test_gen[i]
            predictions = model.predict(batch_x, verbose=0)
            fresh_probs = predictions[1]

            all_fresh_labels.extend(batch_y['fresh_output'])
            all_fresh_probs.extend(fresh_probs.flatten())

        # Calculate ROC curve
        fpr, tpr, thresholds = roc_curve(all_fresh_labels, all_fresh_probs)
        roc_auc = auc(fpr, tpr)

        # Plot ROC curve
        plt.figure(figsize=(10, 8))
        plt.plot(fpr, tpr, color='darkorange', lw=3,
                 label=f'ROC Curve (AUC = {roc_auc:.3f})')
        plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', alpha=0.8)

        plt.xlim([0.0, 1.0])
        plt.ylim([0.0, 1.05])
        plt.xlabel('False Positive Rate', fontsize=14, fontweight='bold')
        plt.ylabel('True Positive Rate', fontsize=14, fontweight='bold')
        plt.title(f'Best Model ({model_name}) - Freshness Classification ROC Curve',
                 fontsize=16, fontweight='bold', pad=20)
        plt.legend(loc="lower right", fontsize=12)
        plt.grid(True, alpha=0.3)
        plt.tight_layout()

        filename = os.path.join(output_dir, 'best_freshness_roc_curve.png')
        plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
        plt.show()

```

```

print(f"✓ Best model ROC curve saved: {filename}")

except Exception as e:
    print(f"Error creating ROC curve for best model: {e}")

def save_best_model_classification_reports(model, test_gen, model_name, fruit_
classes, output_dir="conference_figures"):
    """Save detailed classification reports for best model"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    print(f"Creating classification reports for best model ({model_name})...")

    try:
        # Get predictions
        all_fruit_labels = []
        all_fresh_labels = []
        fruit_pred_classes = []
        fresh_pred_classes = []

        for i in range(len(test_gen)):
            batch_x, batch_y = test_gen[i]
            predictions = model.predict(batch_x, verbose=0)

            all_fruit_labels.extend(np.argmax(batch_y['fruit_output'], axis=
1))
            all_fresh_labels.extend(batch_y['fresh_output'])
            fruit_pred_classes.extend(np.argmax(predictions[0], axis=1))
            fresh_pred_classes.extend((predictions[1] > 0.5).astype(int).flatt
en())

        # Generate classification reports
        fruit_report = classification_report(all_fruit_labels, fruit_pred_clas
ses,
                                           target_names=[f.capitalize() for f
in fruit_classes],
                                           output_dict=True, zero_division=0)
        fresh_report = classification_report(all_fresh_labels, fresh_pred_clas
ses,
                                           target_names=['Fresh', 'Rotten'],
                                           output_dict=True, zero_division=0)

        # Create visualization
        fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 10))
        fig.suptitle(f'Best Model ({model_name}) - Classification Reports', fo
ntsize=18, fontweight='bold')

        # Fruit classification report
        fruit_df = pd.DataFrame(fruit_report).iloc[:-1, :].T
        sns.heatmap(fruit_df.iloc[:, :-1], annot=True, fmt='.3f', cmap='Orange
s',
                    ax=ax1, square=False, linewidths=0.5)
        ax1.set_title('Fruit Classification Report', fontsize=14, fontweight
='bold')
        ax1.set_xlabel('Metrics', fontsize=12, fontweight='bold')
        ax1.set_ylabel('Fruit Classes', fontsize=12, fontweight='bold')

```

```

# Freshness classification report
fresh_df = pd.DataFrame(fresh_report).iloc[: -1, :].T
sns.heatmap(fresh_df.iloc[:, : -1], annot=True, fmt='.3f', cmap='Blue
s',
            ax=ax2, square=False, linewidths=0.5)
ax2.set_title('Freshness Classification Report', fontsize=14, fontweig
ht='bold')
ax2.set_xlabel('Metrics', fontsize=12, fontweight='bold')
ax2.set_ylabel('Freshness Classes', fontsize=12, fontweight='bold')

plt.tight_layout()

filename = os.path.join(output_dir, 'best_classification_reports.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Best model classification reports saved: {filename}")

except Exception as e:
    print(f"Error creating classification reports for best model: {e}")

def save_best_model_confidence_distribution(model, test_gen, model_name, output_dir="conference_figures"):
    """Save prediction confidence distribution for best model"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    print(f"Creating confidence distribution for best model ({model_name})...")

    try:
        # Get prediction confidences
        fruit_confidences = []
        fresh_confidences = []

        for i in range(len(test_gen)):
            batch_x, batch_y = test_gen[i]
            predictions = model.predict(batch_x, verbose=0)

            # Max probability for fruit predictions
            fruit_confidences.extend(np.max(predictions[0], axis=1))
            # Probability for fresh predictions (convert to confidence)
            fresh_probs = predictions[1].flatten()
            fresh_conf = np.maximum(fresh_probs, 1 - fresh_probs)
            fresh_confidences.extend(fresh_conf)

        # Create histograms
        fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))
        fig.suptitle(f'Best Model ({model_name}) - Prediction Confidence Distribution',
                    fontsize=16, fontweight='bold')

        # Fruit confidence histogram
        ax1.hist(fruit_confidences, bins=30, alpha=0.7, color='orange', edgecolor='black')
        ax1.set_title('Fruit Classification Confidence', fontsize=14, fontweight='bold')
        ax1.set_xlabel('Confidence Score', fontsize=12)

```

```

ax1.set_ylabel('Frequency', fontsize=12)
ax1.grid(True, alpha=0.3)
ax1.axvline(np.mean(fruit_confidences), color='red', linestyle='--',
            label=f'Mean: {np.mean(fruit_confidences):.3f}')
ax1.legend()

# Freshness confidence histogram
ax2.hist(fresh_confidences, bins=30, alpha=0.7, color='lightblue', edgecolor='black')
ax2.set_title('Freshness Classification Confidence', fontsize=14, fontweight='bold')
ax2.set_xlabel('Confidence Score', fontsize=12)
ax2.set_ylabel('Frequency', fontsize=12)
ax2.grid(True, alpha=0.3)
ax2.axvline(np.mean(fresh_confidences), color='red', linestyle='--',
            label=f'Mean: {np.mean(fresh_confidences):.3f}')
ax2.legend()

plt.tight_layout()

filename = os.path.join(output_dir, 'best_prediction_confidence_distribution.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Best model confidence distribution saved: {filename}")

except Exception as e:
    print(f"Error creating confidence distribution for best model: {e}")

def save_best_model_sample_predictions(model, test_gen, model_name, fruit_classes, output_dir="conference_figures"):
    """Save sample predictions visualization for best model"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    print(f"Creating sample predictions for best model ({model_name})...")

    try:
        # Get a batch of test data
        batch_x, batch_y = test_gen[0]
        predictions = model.predict(batch_x, verbose=0)

        # Select 12 samples
        num_samples = min(12, len(batch_x))

        fig, axes = plt.subplots(3, 4, figsize=(16, 12))
        fig.suptitle(f'Best Model ({model_name}) - Sample Predictions', fontsize=18, fontweight='bold')

        for i in range(num_samples):
            row, col = i // 4, i % 4

            # Display image
            img = (batch_x[i] * 255).astype(np.uint8)
            axes[row, col].imshow(img)
            axes[row, col].axis('off')

```

```

    # Get predictions
    fruit_pred_idx = np.argmax(predictions[0][i])
    fruit_conf = np.max(predictions[0][i])
    fresh_pred = predictions[1][i][0]
    fresh_class = "Fresh" if fresh_pred > 0.5 else "Rotten"
    fresh_conf = fresh_pred if fresh_pred > 0.5 else 1 - fresh_pred

    # Get actual labels
    actual_fruit_idx = np.argmax(batch_y['fruit_output'][i])
    actual_fresh = "Fresh" if batch_y['fresh_output'][i] > 0.5 else "R
    otten"

    # Create title
    pred_fruit = fruit_classes[fruit_pred_idx].capitalize()
    actual_fruit = fruit_classes[actual_fruit_idx].capitalize()

    title = f"Actual: {actual_fruit}, {actual_fresh}\n"
    title += f"Pred: {pred_fruit} ({fruit_conf:.2f}), {fresh_class}
    ({fresh_conf:.2f})"

    # Color coding: green for correct, red for incorrect
    fruit_correct = fruit_pred_idx == actual_fruit_idx
    fresh_correct = fresh_class == actual_fresh
    color = 'green' if (fruit_correct and fresh_correct) else 'red'

    axes[row, col].set_title(title, fontsize=10, color=color, fontweig
    ht='bold')

plt.tight_layout()

filename = os.path.join(output_dir, 'best_sample_predictions.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Best model sample predictions saved: {filename}")

except Exception as e:
    print(f"Error creating sample predictions for best model: {e}")

def save_vision_transformer_confusion_matrices(model, test_gen, model_name, fr
uit_classes, output_dir="conference_figures"):
    """Save confusion matrices specifically for Vision Transformer model"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    print(f"Creating confusion matrices for {model_name}...")

    try:
        # Get predictions
        all_predictions = []
        all_fruit_labels = []
        all_fresh_labels = []

        for i in range(len(test_gen)):
            batch_x, batch_y = test_gen[i]
            predictions = model.predict(batch_x, verbose=0)
            fruit_preds = predictions[0]
            fresh_preds = predictions[1]

```

```

    all_predictions.append((fruit_preds, fresh_preds))
    all_fruit_labels.extend(np.argmax(batch_y['fruit_output'], axis=
1))

    all_fresh_labels.extend(batch_y['fresh_output'])

# Combine predictions
    fruit_predictions = np.vstack([pred[0] for pred in all_predictions])
    fresh_predictions = np.vstack([pred[1] for pred in all_predictions])
    fruit_pred_classes = np.argmax(fruit_predictions, axis=1)
    fresh_pred_classes = (fresh_predictions > 0.5).astype(int).flatten()

# 1. Fruit Classification Confusion Matrix
    plt.figure(figsize=(12, 10))
    fruit_cm = confusion_matrix(all_fruit_labels, fruit_pred_classes)

# Create proper fruit class names
    display_fruit_classes = [fruit.capitalize() for fruit in fruit_classes]

sns.heatmap(fruit_cm, annot=True, fmt='d', cmap='Oranges',
            xticklabels=display_fruit_classes,
            yticklabels=display_fruit_classes,
            square=True, linewidths=0.5)
plt.title(f'{model_name} - Fruit Classification Confusion Matrix',
        fontsize=16, fontweight='bold', pad=20)
plt.xlabel('Predicted Fruit Class', fontsize=14, fontweight='bold')
plt.ylabel('Actual Fruit Class', fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()

    safe_model_name = model_name.lower().replace(' ', '_').replace('+',
'_').replace('-', '_')
    filename = os.path.join(output_dir, f'{safe_model_name}_fruit_confusion_matrix.png')
    plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
    plt.show()
    print(f"✓ {model_name} fruit confusion matrix saved: {filename}")

# 2. Freshness Classification Confusion Matrix
    plt.figure(figsize=(8, 6))
    fresh_cm = confusion_matrix(all_fresh_labels, fresh_pred_classes)

    sns.heatmap(fresh_cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Fresh', 'Rotten'],
            yticklabels=['Fresh', 'Rotten'],
            square=True, linewidths=0.5)
    plt.title(f'{model_name} - Freshness Classification Confusion Matrix',
            fontsize=16, fontweight='bold', pad=20)
    plt.xlabel('Predicted Freshness Class', fontsize=14, fontweight='bold')

    plt.ylabel('Actual Freshness Class', fontsize=14, fontweight='bold')
    plt.tight_layout()

    filename = os.path.join(output_dir, f'{safe_model_name}_freshness_confusion_matrix.png')

```

```

plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ {model_name} freshness confusion matrix saved: {filename}")

except Exception as e:
    print(f"Error creating confusion matrices for {model_name}: {e}")

def save_results_summary_visualization(results_df, output_dir="conference_figures"):
    """Save a comprehensive results summary visualization"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    # Create a comprehensive summary plot
    fig, axes = plt.subplots(2, 2, figsize=(20, 16))
    fig.suptitle('Multi-Task Learning Results Summary', fontsize=20, fontweight='bold')

    # 1. Fruit Classification Accuracy Comparison
    models = results_df['Model']
    x_pos = np.arange(len(models))
    # FIXED: Updated colors for 9 models
    colors = ['#FF6B6B', '#4ECDC4', '#45B7D1', '#96CEB4', '#FECA57', '#FF9FF3', '#A8E6CF', '#FFD93D', '#FF8A80']

    bars1 = axes[0,0].bar(x_pos, results_df['Test Fruit Acc'],
                        color=colors[:len(models)], alpha=0.8, edgecolor='black', linewidth=1)
    axes[0,0].set_title('Fruit Classification Test Accuracy', fontsize=16, fontweight='bold')
    axes[0,0].set_xlabel('Models', fontsize=12)
    axes[0,0].set_ylabel('Accuracy', fontsize=12)
    axes[0,0].set_xticks(x_pos)
    axes[0,0].set_xticklabels(models, rotation=45, ha='right')
    axes[0,0].grid(True, alpha=0.3, axis='y')

    # Add value labels
    for bar in bars1:
        height = bar.get_height()
        axes[0,0].text(bar.get_x() + bar.get_width()/2., height + 0.01,
                      f'{height:.3f}', ha='center', va='bottom', fontweight='bold')

    # 2. Freshness Classification Accuracy Comparison
    bars2 = axes[0,1].bar(x_pos, results_df['Test Fresh Acc'],
                        color=colors[:len(models)], alpha=0.8, edgecolor='black', linewidth=1)
    axes[0,1].set_title('Freshness Classification Test Accuracy', fontsize=16, fontweight='bold')
    axes[0,1].set_xlabel('Models', fontsize=12)
    axes[0,1].set_ylabel('Accuracy', fontsize=12)
    axes[0,1].set_xticks(x_pos)
    axes[0,1].set_xticklabels(models, rotation=45, ha='right')
    axes[0,1].grid(True, alpha=0.3, axis='y')

    for bar in bars2:
        height = bar.get_height()

```

```

        axes[0,1].text(bar.get_x() + bar.get_width()/2., height + 0.01,
                       f'{height:.3f}', ha='center', va='bottom', fontweight='bold')

    # 3. Combined Performance Comparison
    axes[1,0].plot(x_pos, results_df['Test Fruit Acc'], 'o-', label='Fruit Classification',
                  linewidth=3, markersize=10, color='#FF6B6B')
    axes[1,0].plot(x_pos, results_df['Test Fresh Acc'], 's-', label='Freshness Classification',
                  linewidth=3, markersize=10, color='#4ECDC4')
    axes[1,0].set_title('Multi-Task Performance Comparison', fontsize=16, fontweight='bold')
    axes[1,0].set_xlabel('Models', fontsize=12)
    axes[1,0].set_ylabel('Accuracy', fontsize=12)
    axes[1,0].set_xticks(x_pos)
    axes[1,0].set_xticklabels(models, rotation=45, ha='right')
    axes[1,0].legend(fontsize=12)
    axes[1,0].grid(True, alpha=0.3)

    # 4. Loss Comparison
    axes[1,1].plot(x_pos, results_df['Test Loss'], 'o-', label='Test Loss',
                  linewidth=3, markersize=10, color='#45B7D1')
    axes[1,1].set_title('Model Loss Comparison', fontsize=16, fontweight='bold')
    axes[1,1].set_xlabel('Models', fontsize=12)
    axes[1,1].set_ylabel('Loss', fontsize=12)
    axes[1,1].set_xticks(x_pos)
    axes[1,1].set_xticklabels(models, rotation=45, ha='right')
    axes[1,1].legend(fontsize=12)
    axes[1,1].grid(True, alpha=0.3)

plt.tight_layout()

filename = os.path.join(output_dir, 'multi_task_results_summary.png')
plt.savefig(filename, dpi=300, bbox_inches='tight', facecolor='white')
plt.show()
print(f"✓ Results summary visualization saved: {filename}")

def save_all_individual_visualizations(results_df, training_histories, best_model_name, fruit_classes,
                                     best_model=None, test_gen=None, trained_models=None, output_dir="conference_figures"):
    """Save all individual visualizations for conference paper - FIXED FOR 9 MODELS"""
    print(f"\n{'='*70}")
    print("CREATING CONFERENCE-READY VISUALIZATIONS")
    print(f"{'='*70}")

    # Create output directory
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)
        print(f"Created output directory: {output_dir}")

    # 1. Individual accuracy comparisons
    save_individual_fruit_accuracy_comparison(results_df, output_dir)
    save_individual_freshness_accuracy_comparison(results_df, output_dir)

```

```

# 2. Combined performance plots
save_individual_combined_performance_plot(results_df, output_dir)
save_individual_loss_comparison(results_df, output_dir)

# 3. Performance heatmaps
save_individual_performance_heatmaps(results_df, output_dir)

# 4. Correlation and average performance
save_individual_correlation_scatter(results_df, output_dir)
save_individual_average_performance_bar(results_df, output_dir)

# 5. Training history for each model
for model_name, history in training_histories.items():
    save_individual_training_history(history, model_name, output_dir)

# 6. Training history for best model
if best_model_name in training_histories:
    save_best_model_training_history(training_histories[best_model_name],
best_model_name, output_dir)

# 7. Best model detailed analysis
if best_model is not None and test_gen is not None:
    save_best_model_confusion_matrices(best_model, test_gen, best_model_name, fruit_classes, output_dir)
    save_best_model_roc_curve(best_model, test_gen, best_model_name, output_dir)
    save_best_model_classification_reports(best_model, test_gen, best_model_name, fruit_classes, output_dir)
    save_best_model_confidence_distribution(best_model, test_gen, best_model_name, output_dir)
    save_best_model_sample_predictions(best_model, test_gen, best_model_name, fruit_classes, output_dir)

# 8. ViT-Base Pretrained confusion matrices (FIXED)
if trained_models is not None:
    vit_model_for_viz = trained_models.get("ViT-Base Pretrained")
    if vit_model_for_viz is not None and test_gen is not None:
        save_vision_transformer_confusion_matrices(vit_model_for_viz, test_gen, "ViT-Base Pretrained", fruit_classes, output_dir)

# 9. Save the comprehensive results summary
save_results_summary_visualization(results_df, output_dir)

print(f"\n{'='*70}")
print("ALL VISUALIZATIONS SAVED SUCCESSFULLY!")
print(f"Check the '{output_dir}' folder for all conference-ready figures")
print(f"{'='*70}")

#=====
# MAIN EXECUTION FUNCTION - FIXED VERSION
#=====

def main():
    """Main execution function for enhanced multi-task Learning - FIXED VERSION"""
    print("Starting Enhanced Fresh vs Rotten Fruit Multi-Task Classification")

```

```

print("="*80)

try:
    # Step 1: Load and preprocess data
    df_raw = load_and_process_dataset()

    if len(df_raw) == 0:
        print("⚠️ No data loaded, exiting...")
        return None, None, None, None, None

    # Step 2: Apply smart balancing
    df_balanced = apply_smart_balancing(df_raw)

    if len(df_balanced) == 0:
        print("⚠️ No data after balancing, exiting...")
        return None, None, None, None, None

    # Step 3: Create stratified splits
    train_df, val_df, test_df = create_stratified_splits(df_balanced)

    if len(train_df) == 0 or len(val_df) == 0 or len(test_df) == 0:
        print("⚠️ Empty splits created, exiting...")
        return None, None, None, None, None

    # Step 4: Prepare labels and classes
    fruit_classes = sorted(df_balanced['fruit'].unique())
    num_fruit_classes = len(fruit_classes)

    if num_fruit_classes == 0:
        print("⚠️ No fruit classes found, exiting...")
        return None, None, None, None, None

    print(f"\n🎯 Target fruit classes ({num_fruit_classes}): {fruit_classes}")

    # Step 5: Compute class weights
    fruit_class_weights, fresh_class_weights = compute_class_weights(train_df)

    # Step 6: Create data generators
    try:
        train_datagen = ImageDataGenerator(
            rescale=1./255,
            rotation_range=20,
            width_shift_range=0.15,
            height_shift_range=0.15,
            shear_range=0.15,
            zoom_range=0.15,
            horizontal_flip=True,
            brightness_range=[0.9, 1.1],
            fill_mode='nearest'
        )

        val_test_datagen = ImageDataGenerator(rescale=1./255)

        batch_size = min(4, len(train_df) // 20) # FIXED: Reduced for memory efficiency

```

```

batch_size = max(1, batch_size) # Ensure at least 1

train_generator = EnhancedMultiTaskDataGenerator(
    train_df, train_datagen, target_size=(224, 224), batch_size=batch_size,
    fruit_classes=fruit_classes, shuffle=True
)
val_generator = EnhancedMultiTaskDataGenerator(
    val_df, val_test_datagen, target_size=(224, 224), batch_size=batch_size,
    fruit_classes=fruit_classes, shuffle=False
)
test_generator = EnhancedMultiTaskDataGenerator(
    test_df, val_test_datagen, target_size=(224, 224), batch_size=batch_size,
    fruit_classes=fruit_classes, shuffle=False
)

print(f"\n📊 Data Generator Summary:")
print(f"Training samples: {len(train_df)} (batches: {len(train_generator)})")
print(f"Validation samples: {len(val_df)} (batches: {len(val_generator)})")
print(f"Test samples: {len(test_df)} (batches: {len(test_generator)})")
print(f"Batch size: {batch_size}")

except Exception as e:
    print(f"⚠️ Error creating data generators: {e}")
    return None, None, None, None

# Step 7: Define models to train
# Memory cleanup function
def cleanup_memory():
    """Clean up memory between model training"""
    import gc
    try:
        gc.collect()
        tf.keras.backend.clear_session()
    except:
        pass

models_config = [
    ("Custom CNN", lambda: create_multi_task_custom_cnn(num_fruit_classes=num_fruit_classes)),
    ("VGG16", lambda: create_multi_task_vgg16(num_fruit_classes=num_fruit_classes)),
    ("InceptionV3", lambda: create_multi_task_inceptionv3(num_fruit_classes=num_fruit_classes)),
    ("ResNet152", lambda: create_multi_task_resnet152(num_fruit_classes=num_fruit_classes)),
    ("YOLOv3+FC", lambda: create_multi_task_yolov3_backbone(num_fruit_classes=num_fruit_classes)),
    ("EfficientNetB4", lambda: create_multi_task_efficientnet(num_fruit_classes=num_fruit_classes)),
    ("DenseNet121", lambda: create_multi_task_densenet(num_fruit_classes=num_fruit_classes))
]

```

```

es=num_fruit_classes)),
    ("ViT-Base Pretrained", lambda: create_multi_task_vit_base_pretrained(num_fruit_classes=num_fruit_classes)),
    ("ConvNeXt-Base", lambda: create_multi_task_convnext_base(num_fruit_classes=num_fruit_classes))
]

# Storage for results
all_results = []
trained_models = {}
training_histories = {}

# Step 8: Train models
for model_name, model_creator in models_config:
    try:
        print(f"\n{'='*60}")
        print(f"Creating and training {model_name}")
        print(f"{'='*60}")

        # Create model
        model = model_creator()
        print(f"✅ {model_name} created successfully!")

        try:
            param_count = model.count_params()
            print(f"Total parameters: {param_count:,}")
        except:
            print("Parameter count not available")

        # Train model
        epochs = 10 if model_name == "ViT-Base Pretrained" else 8 # Reduced epochs for faster training
        trained_model, history = compile_and_train_multi_task_model(
            model, model_name, train_generator, val_generator,
            epochs=epochs, alpha=0.7
        )

        # Evaluate model
        results = evaluate_multi_task_model(trained_model, train_generator, val_generator, test_generator, model_name)
        all_results.append(results)

        # Store model and history
        trained_models[model_name] = trained_model
        training_histories[model_name] = history

        print(f"✅ {model_name} Results:")
        print(f"    Test Fruit Acc: {results['Test Fruit Acc']:.4f}")
        print(f"    Test Fresh Acc: {results['Test Fresh Acc']:.4f}")

        # Clear memory
        cleanup_memory()

    except Exception as e:
        print(f"❌ Error training {model_name}: {str(e)}")
        # Add dummy results
        dummy_results = {

```

```

        'Model': model_name,
        'Train Loss': 999.0, 'Valid Loss': 999.0, 'Test Loss': 99
9.0,
        'Train Fruit Acc': 0.0, 'Valid Fruit Acc': 0.0, 'Test Fruit Acc': 0.0,
        'Train Fresh Acc': 0.0, 'Valid Fresh Acc': 0.0, 'Test Fresh Acc': 0.0
    }
    all_results.append(dummy_results)
    cleanup_memory()
    continue

# Step 9: Create results DataFrame
if len(all_results) == 0:
    print("⚠️ No results to process")
    return None, None, None, None, None

results_df = pd.DataFrame(all_results)

# Calculate combined accuracy and find best model
results_df['Combined_Test_Acc'] = (results_df['Test Fruit Acc'] + results_df['Test Fresh Acc']) / 2

if len(results_df) > 0:
    best_model_idx = results_df['Combined_Test_Acc'].idxmax()
    best_model_name = results_df.loc[best_model_idx, 'Model']
    best_model = trained_models.get(best_model_name)
else:
    best_model_name = "None"
    best_model = None

print(f"\n{'='*80}")
print(f"🏆 BEST MODEL: {best_model_name}")
if best_model_idx is not None:
    print(f"Test Fruit Accuracy: {results_df.loc[best_model_idx, 'Test Fruit Acc']:.4f}")
    print(f"Test Fresh Accuracy: {results_df.loc[best_model_idx, 'Test Fresh Acc']:.4f}")
    print(f"Combined Test Accuracy: {results_df.loc[best_model_idx, 'Combined_Test_Acc']:.4f}")
    print(f"\n{'='*80}")

# Step 10: Display summary table
print("\n" + "="*120)
print("FINAL MULTI-TASK LEARNING SUMMARY TABLE")
print("="*120)

summary_df = results_df.copy()
for col in ['Train Loss', 'Valid Loss', 'Test Loss', 'Combined_Test_Acc']:
    summary_df[col] = summary_df[col].apply(lambda x: f"{x:.4f}")
for col in ['Train Fruit Acc', 'Valid Fruit Acc', 'Test Fruit Acc', 'Train Fresh Acc', 'Valid Fresh Acc', 'Test Fresh Acc']:
    summary_df[col] = summary_df[col].apply(lambda x: f"{x:.4f}")

print(summary_df.to_string(index=False))

```

```

# Step 11: Save results
try:
    results_df.to_csv('multi_task_model_comparison_results.csv', index
=False)
    print(f"\n✅ Results saved to 'multi_task_model_comparison_result
s.csv'")
except Exception as e:
    print(f"⚠️ Error saving results CSV: {e}")

# Step 12: Create ALL visualizations
try:
    save_all_individual_visualizations(
        results_df, training_histories, best_model_name, fruit_classes,
        best_model, test_generator, trained_models
    )
except Exception as e:
    print(f"⚠️ Error creating visualizations: {e}")

print(f"\n{'='*80}")
print("✅ ENHANCED MULTI-TASK LEARNING PROJECT COMPLETED!")
print(f"🏆 Best Model: {best_model_name}")
print(f"🎯 Target fruit classes: {fruit_classes}")
print(f"📊 Total models trained: {len(results_df)}")
print(f"📁 All visualizations saved in 'conference_figures' folder")
print(f"{'='*80}")

return results_df, trained_models, best_model_name, fruit_classes, test_generator

except Exception as e:
    print(f"❌ Fatal error in main function: {e}")
    import traceback
    traceback.print_exc()
    return None, None, None, None, None

# Execute the main function with proper error handling
if __name__ == "__main__":
    try:
        results_df, trained_models, best_model_name, fruit_classes, test_gen =
main()

        if results_df is not None:
            print("\n" + "="*80)
            print("🎉 PROJECT COMPLETED SUCCESSFULLY! 🎉")
            print("="*80)
            print(f"📊 Results available in DataFrame with {len(results_df)}
models")
            print(f"🏆 Best performing model: {best_model_name}")
            if len(results_df) > 0:
                print(f"📈 Best combined accuracy: {results_df['Combined_Test
_Acc'].max():.4f}")
                print(f"🍎 Best fruit accuracy: {results_df['Test Fruit Acc
c'].max():.4f}")
                print(f"🌿 Best fresh accuracy: {results_df['Test Fresh Acc
c'].max():.4f}")
                print(f"🎯 Fruit classes handled: {len(fruit_classes) if fruit_cl

```

```

asses else 0}"))

    # Print detailed results
    print(f"\n📄 DETAILED RESULTS:")
    print("="*50)
    for _, row in results_df.iterrows():
        print(f"• {row['Model']}:")
        print(f"    🍏 Fruit: {row['Test Fruit Acc']:.4f}, Fresh: {row['T
est Fresh Acc']:.4f}, Combined: {row['Combined_Test_Acc']:.4f}")

    print(f"\n📁 FILES GENERATED:")
    print("="*30)
    print(" • multi_task_model_comparison_results.csv - Main results")
    print(" • conference_figures/ - All visualization files")

    print(f"\n🎨 VISUALIZATIONS CREATED:")
    print("="*35)
    print(" • Individual accuracy comparisons")
    print(" • Combined performance plots")
    print(" • Loss comparisons")
    print(" • Performance heatmaps")
    print(" • Correlation scatter plots")
    print(" • Average performance bars")
    print(" • Training history for each model")
    print(" • Best model training history")
    print(" • Best model confusion matrices")
    print(" • Best model ROC curves")
    print(" • Best model classification reports")
    print(" • Best model confidence distributions")
    print(" • Best model sample predictions")
    print(" • ViT-Base Pretrained confusion matrices")
    print(" • Comprehensive results summary")

    print("="*80)
    else:
        print("\n" + "="*80)
        print("⚠️ PROJECT COMPLETED WITH ISSUES")
        print("Some errors occurred during execution. Check the logs abov
e.")

        print("="*80)

    except Exception as e:
        print(f"\n❌ CRITICAL ERROR: {e}")
        import traceback
        traceback.print_exc()
        print("\nThe program encountered a critical error and could not comple
te.")

    print("\n" + "="*80)
    print("🎉 ENHANCED MULTI-TASK FRUIT CLASSIFICATION PROJECT")
    print("🎉 WITH ALL ISSUES FIXED FOR 9 MODELS - COMPLETED!")
    print("="*80)

```

```
2025-08-09 06:40:43.545498: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1754721643.747824      36 cuda_dnn.cc:8310] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1754721643.811401      36 cuda_blas.cc:1418] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
```

GPU available: 1 GPU(s)

TensorFlow version: 2.18.0

Keras version: 3.8.0

Starting Enhanced Fresh vs Rotten Fruit Multi-Task Classification

=====

===

 Starting Enhanced Dataset Loading and Processing...

Processing Train: 100%|██████████| 18/18 [01:12<00:00, 4.04s/it]

Processing Test: 100%|██████████| 14/14 [00:18<00:00, 1.29s/it]

✓ Total images loaded: 30357

📊 Fruit distribution:

fruit

apple 7451

banana 7192

tomato 4291

orange 3852

potato 1978

capsicum 1891

okra 1567

cucumber 1451

bittergourd 684

Name: count, dtype: int64

📄 Filtered to target fruits: 30357 images

🎯 Target fruit distribution:

fruit

apple 7451

banana 7192

tomato 4291

orange 3852

potato 1978

capsicum 1891

okra 1567

cucumber 1451

bittergourd 684

Name: count, dtype: int64

🌸 Applying Smart Balancing Strategy...

🍌 Processing banana: 7192 samples

✓ Sampled 3000 images (Fresh: 1401, Rotten: 1599)

🍌 Processing potato: 1978 samples

✓ Sampled 1978 images (Fresh: 806, Rotten: 1172)

🍌 Processing okra: 1567 samples

✓ Sampled 1567 images (Fresh: 1005, Rotten: 562)

🍌 Processing tomato: 4291 samples

✓ Sampled 3000 images (Fresh: 1477, Rotten: 1523)

🍌 Processing orange: 3852 samples

✓ Sampled 3000 images (Fresh: 1443, Rotten: 1557)

🍌 Processing cucumber: 1451 samples

✓ Sampled 1451 images (Fresh: 775, Rotten: 676)

🍌 Processing capsicum: 1891 samples

✓ Sampled 1891 images (Fresh: 990, Rotten: 901)

🍌 Processing apple: 7451 samples

✓ Sampled 3000 images (Fresh: 1294, Rotten: 1706)

🍌 Processing bittergourd: 684 samples

🔄 Need to augment 116 additional images

✓ Created 1368 samples (684 original + 684 augmented)

Final Balanced Dataset Statistics:

Total samples: 20255

Fruit distribution:

fruit

banana	3000
orange	3000
tomato	3000
apple	3000
potato	1978
capsicum	1891
okra	1567
cucumber	1451
bittergourd	1368

Name: count, dtype: int64

Freshness distribution:

fresh

0	10410
1	9845

Name: count, dtype: int64

Augmentation Statistics:

bittergourd: 684 → 1368 (+684 augmented)

Creating Stratified Train/Val/Test Splits (70/15/15)...

✓ Train: 14178 (70.0%)

✓ Val: 3038 (15.0%)

✓ Test: 3039 (15.0%)

🎯 Target fruit classes (9): ['apple', 'banana', 'bittergourd', 'capsicum', 'cucumber', 'okra', 'orange', 'potato', 'tomato']

Computing Class Weights...

🍎 Fruit class weights: {0: 0.7501587301587301, 1: 0.7501587301587301, 2: 1.6443980514961727, 3: 1.189828801611279, 4: 1.5505249343832022, 5: 1.4373479318734794, 6: 0.7501587301587301, 7: 1.1382466281310213, 8: 0.7501587301587301}

🌱 Fresh class weights: {0: 0.9729618446335437, 1: 1.0285838653511317}

Data Generator Summary:

Training samples: 14178 (batches: 3545)

Validation samples: 3038 (batches: 760)

Test samples: 3039 (batches: 760)

Batch size: 4

=====
Creating and training Custom CNN
=====

I0000 00:00:1754721747.996343 36 gpu_device.cc:2022] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 15513 MB memory: -> device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute capability: 6.0

✓ Custom CNN created successfully!

Total parameters: 27,048,746

=====

Training Custom CNN

=====

Epoch 1/8

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1754721758.760866 104 service.cc:148] XLA service 0x7a0f3003fa0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1754721758.761766 104 service.cc:156] StreamExecutor device (0): Tesla P100-PCIE-16GB, Compute Capability 6.0

I0000 00:00:1754721759.689415 104 cuda_dnn.cc:529] Loaded cuDNN version 90300

4/3545 ----- 2:45 47ms/step - fresh_output_accuracy: 0.5990
- fresh_output_loss: 0.8430 - fruit_output_accuracy: 0.2083 - fruit_output_loss: 2.8324 - loss: 2.2356

I0000 00:00:1754721768.628224 104 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

3545/3545 ————— **218s** 56ms/step - fresh_output_accuracy: 0.5246
 - fresh_output_loss: 0.7157 - fruit_output_accuracy: 0.3759 - fruit_output_loss: 1.6297 - loss: 1.3554 - val_fresh_output_accuracy: 0.5227 - val_fresh_output_loss: 1.3180 - val_fruit_output_accuracy: 0.5201 - val_fruit_output_loss: 2.4397 - val_loss: 2.1040 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **89s** 25ms/step - fresh_output_accuracy: 0.6174
 - fresh_output_loss: 0.6535 - fruit_output_accuracy: 0.5636 - fruit_output_loss: 1.0880 - loss: 0.9576 - val_fresh_output_accuracy: 0.6649 - val_fresh_output_loss: 0.5928 - val_fruit_output_accuracy: 0.6458 - val_fruit_output_loss: 0.8495 - val_loss: 0.7727 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **89s** 25ms/step - fresh_output_accuracy: 0.6503
 - fresh_output_loss: 0.6204 - fruit_output_accuracy: 0.5994 - fruit_output_loss: 0.9830 - loss: 0.8743 - val_fresh_output_accuracy: 0.6577 - val_fresh_output_loss: 0.8041 - val_fruit_output_accuracy: 0.6485 - val_fruit_output_loss: 1.0793 - val_loss: 0.9970 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **89s** 25ms/step - fresh_output_accuracy: 0.6518
 - fresh_output_loss: 0.6090 - fruit_output_accuracy: 0.6281 - fruit_output_loss: 0.9203 - loss: 0.8267 - val_fresh_output_accuracy: 0.6504 - val_fresh_output_loss: 0.6025 - val_fruit_output_accuracy: 0.6771 - val_fruit_output_loss: 0.7848 - val_loss: 0.7303 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **88s** 25ms/step - fresh_output_accuracy: 0.6568
 - fresh_output_loss: 0.5984 - fruit_output_accuracy: 0.6374 - fruit_output_loss: 0.8725 - loss: 0.7902 - val_fresh_output_accuracy: 0.6764 - val_fresh_output_loss: 0.5703 - val_fruit_output_accuracy: 0.7110 - val_fruit_output_loss: 0.6700 - val_loss: 0.6403 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **88s** 25ms/step - fresh_output_accuracy: 0.6791
 - fresh_output_loss: 0.5810 - fruit_output_accuracy: 0.6713 - fruit_output_loss: 0.7806 - loss: 0.7208 - val_fresh_output_accuracy: 0.7488 - val_fresh_output_loss: 0.5509 - val_fruit_output_accuracy: 0.7656 - val_fruit_output_loss: 1.0155 - val_loss: 0.8766 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **87s** 25ms/step - fresh_output_accuracy: 0.6911
 - fresh_output_loss: 0.5585 - fruit_output_accuracy: 0.6992 - fruit_output_loss: 0.7354 - loss: 0.6823 - val_fresh_output_accuracy: 0.7380 - val_fresh_output_loss: 0.5620 - val_fruit_output_accuracy: 0.7755 - val_fruit_output_loss: 1.1375 - val_loss: 0.9652 - learning_rate: 0.0010

Epoch 8/8

3543/3545 ————— **0s** 21ms/step - fresh_output_accuracy: 0.6949 - fresh_output_loss: 0.5616 - fruit_output_accuracy: 0.7176 - fruit_output_loss: 0.6830 - loss: 0.6466

Epoch 8: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

3545/3545 ————— **86s** 24ms/step - fresh_output_accuracy: 0.6949
 - fresh_output_loss: 0.5616 - fruit_output_accuracy: 0.7176 - fruit_output_loss: 0.6830 - loss: 0.6466 - val_fresh_output_accuracy: 0.7633 - val_fresh_output_loss: 0.4922 - val_fruit_output_accuracy: 0.8140 - val_fruit_output_loss: 1.0861 - val_loss: 0.9084 - learning_rate: 0.0010

Restoring model weights from the end of the best epoch: 5.

Evaluating Custom CNN...

✓ Custom CNN Results:
 Test Fruit Acc: 0.7154
 Test Fresh Acc: 0.6808

=====
Creating and training VGG16

=====
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5

58889256/58889256 ————— **0s** 0us/step

✓ VGG16 created successfully!

Total parameters: 15,161,290

=====
Training VGG16

=====
Epoch 1/8

3545/3545 ————— **120s** 31ms/step - fresh_output_accuracy: 0.5121
- fresh_output_loss: 0.7139 - fruit_output_accuracy: 0.2649 - fruit_output_loss: 1.9291 - loss: 1.5645 - val_fresh_output_accuracy: 0.5342 - val_fresh_output_loss: 0.6928 - val_fruit_output_accuracy: 0.4809 - val_fruit_output_loss: 1.9377 - val_loss: 1.5635 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **104s** 29ms/step - fresh_output_accuracy: 0.5282
- fresh_output_loss: 0.6920 - fruit_output_accuracy: 0.4725 - fruit_output_loss: 1.2774 - loss: 1.1014 - val_fresh_output_accuracy: 0.5135 - val_fresh_output_loss: 0.7174 - val_fruit_output_accuracy: 0.3657 - val_fruit_output_loss: 2.4938 - val_loss: 1.9595 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **105s** 30ms/step - fresh_output_accuracy: 0.5279
- fresh_output_loss: 0.6907 - fruit_output_accuracy: 0.6204 - fruit_output_loss: 0.9492 - loss: 0.8716 - val_fresh_output_accuracy: 0.5543 - val_fresh_output_loss: 0.6841 - val_fruit_output_accuracy: 0.7146 - val_fruit_output_loss: 0.6966 - val_loss: 0.6923 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **105s** 30ms/step - fresh_output_accuracy: 0.5417
- fresh_output_loss: 0.6879 - fruit_output_accuracy: 0.6772 - fruit_output_loss: 0.8557 - loss: 0.8054 - val_fresh_output_accuracy: 0.5592 - val_fresh_output_loss: 0.6865 - val_fruit_output_accuracy: 0.8585 - val_fruit_output_loss: 0.3180 - val_loss: 0.4287 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **105s** 30ms/step - fresh_output_accuracy: 0.5505
- fresh_output_loss: 0.6861 - fruit_output_accuracy: 0.7639 - fruit_output_loss: 0.6701 - loss: 0.6749 - val_fresh_output_accuracy: 0.5708 - val_fresh_output_loss: 1.6829 - val_fruit_output_accuracy: 0.7959 - val_fruit_output_loss: 45.4895 - val_loss: 32.3676 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **105s** 30ms/step - fresh_output_accuracy: 0.5554
- fresh_output_loss: 0.6844 - fruit_output_accuracy: 0.7962 - fruit_output_loss: 0.5512 - loss: 0.5912 - val_fresh_output_accuracy: 0.5642 - val_fresh_output_loss: 0.6729 - val_fruit_output_accuracy: 0.9236 - val_fruit_output_loss: 0.2019 - val_loss: 0.3433 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **105s** 30ms/step - fresh_output_accuracy: 0.5549
- fresh_output_loss: 0.6841 - fruit_output_accuracy: 0.8299 - fruit_output_loss: 0.4632 - loss: 0.5295 - val_fresh_output_accuracy: 0.5820 - val_fresh_output_loss: 0.7053 - val_fruit_output_accuracy: 0.8598 - val_fruit_output_loss: 19.1502 - val_loss: 13.6255 - learning_rate: 0.0010

Epoch 8/8

3545/3545 ————— **104s** 29ms/step - fresh_output_accuracy: 0.5689
- fresh_output_loss: 0.6796 - fruit_output_accuracy: 0.8513 - fruit_output_loss:

ss: 0.3960 - loss: 0.4811 - val_fresh_output_accuracy: 0.6248 - val_fresh_output_loss: 0.7527 - val_fruit_output_accuracy: 0.9467 - val_fruit_output_loss: 15.1404 - val_loss: 10.8311 - learning_rate: 0.0010
 Restoring model weights from the end of the best epoch: 6.
 Evaluating VGG16...

✅ VGG16 Results:
 Test Fruit Acc: 0.9296
 Test Fresh Acc: 0.5676

=====
 Creating and training InceptionV3
 =====

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
87910968/87910968 ————— **0s** 0us/step

✅ InceptionV3 created successfully!
 Total parameters: 23,035,818

=====
 Training InceptionV3
 =====

Epoch 1/8

3545/3545 ————— **138s** 32ms/step - fresh_output_accuracy: 0.7605 - fresh_output_loss: 0.4955 - fruit_output_accuracy: 0.7742 - fruit_output_loss: 0.6937 - loss: 0.6342 - val_fresh_output_accuracy: 0.9161 - val_fresh_output_loss: 0.2191 - val_fruit_output_accuracy: 0.9700 - val_fruit_output_loss: 0.1039 - val_loss: 0.1385 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **92s** 26ms/step - fresh_output_accuracy: 0.8561 - fresh_output_loss: 0.3307 - fruit_output_accuracy: 0.9077 - fruit_output_loss: 0.2962 - loss: 0.3065 - val_fresh_output_accuracy: 0.9269 - val_fresh_output_loss: 0.1714 - val_fruit_output_accuracy: 0.9793 - val_fruit_output_loss: 0.0659 - val_loss: 0.0976 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.8865 - fresh_output_loss: 0.2631 - fruit_output_accuracy: 0.9342 - fruit_output_loss: 0.2107 - loss: 0.2263 - val_fresh_output_accuracy: 0.9296 - val_fresh_output_loss: 0.1722 - val_fruit_output_accuracy: 0.9819 - val_fruit_output_loss: 0.0680 - val_loss: 0.0993 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.8949 - fresh_output_loss: 0.2493 - fruit_output_accuracy: 0.9442 - fruit_output_loss: 0.1964 - loss: 0.2123 - val_fresh_output_accuracy: 0.9414 - val_fresh_output_loss: 0.1375 - val_fruit_output_accuracy: 0.9901 - val_fruit_output_loss: 0.0417 - val_loss: 0.0704 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.9059 - fresh_output_loss: 0.2213 - fruit_output_accuracy: 0.9489 - fruit_output_loss: 0.1638 - loss: 0.1810 - val_fresh_output_accuracy: 0.9463 - val_fresh_output_loss: 0.1443 - val_fruit_output_accuracy: 0.9878 - val_fruit_output_loss: 0.0442 - val_loss: 0.0743 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **90s** 25ms/step - fresh_output_accuracy: 0.9097 - fresh_output_loss: 0.2120 - fruit_output_accuracy: 0.9581 - fruit_output_loss: 0.1376 - loss: 0.1599 - val_fresh_output_accuracy: 0.9434 - val_fresh_output_loss: 0.1374 - val_fruit_output_accuracy: 0.9908 - val_fruit_output_loss: 0.0394 - val_loss: 0.0689 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.9198
 - fresh_output_loss: 0.1925 - fruit_output_accuracy: 0.9541 - fruit_output_loss: 0.1620 - loss: 0.1711 - val_fresh_output_accuracy: 0.9510 - val_fresh_output_loss: 0.1077 - val_fruit_output_accuracy: 0.9934 - val_fruit_output_loss: 0.0325 - val_loss: 0.0551 - learning_rate: 0.0010

Epoch 8/8

3545/3545 ————— **92s** 26ms/step - fresh_output_accuracy: 0.9280
 - fresh_output_loss: 0.1830 - fruit_output_accuracy: 0.9618 - fruit_output_loss: 0.1281 - loss: 0.1446 - val_fresh_output_accuracy: 0.9536 - val_fresh_output_loss: 0.1215 - val_fruit_output_accuracy: 0.9908 - val_fruit_output_loss: 0.0400 - val_loss: 0.0645 - learning_rate: 0.0010

Restoring model weights from the end of the best epoch: 7.

Evaluating InceptionV3...

✓ InceptionV3 Results:
 Test Fruit Acc: 0.9908
 Test Fresh Acc: 0.9552

=====
 Creating and training ResNet152
 =====

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet152v2_weights_tf_dim_ordering_tf_kernels_notop.h5

234545216/234545216 ————— **1s** 0us/step

✓ ResNet152 created successfully!

Total parameters: 59,564,682

=====
 Training ResNet152
 =====

Epoch 1/8

3545/3545 ————— **250s** 58ms/step - fresh_output_accuracy: 0.7194
 - fresh_output_loss: 0.5493 - fruit_output_accuracy: 0.7442 - fruit_output_loss: 0.7898 - loss: 0.7177 - val_fresh_output_accuracy: 0.9184 - val_fresh_output_loss: 0.1852 - val_fruit_output_accuracy: 0.9710 - val_fruit_output_loss: 0.0899 - val_loss: 0.1186 - learning_rate: 5.0000e-04

Epoch 2/8

3545/3545 ————— **178s** 50ms/step - fresh_output_accuracy: 0.8556
 - fresh_output_loss: 0.3203 - fruit_output_accuracy: 0.9341 - fruit_output_loss: 0.2215 - loss: 0.2512 - val_fresh_output_accuracy: 0.9332 - val_fresh_output_loss: 0.1430 - val_fruit_output_accuracy: 0.9865 - val_fruit_output_loss: 0.0386 - val_loss: 0.0700 - learning_rate: 5.0000e-04

Epoch 3/8

3545/3545 ————— **184s** 52ms/step - fresh_output_accuracy: 0.8893
 - fresh_output_loss: 0.2486 - fruit_output_accuracy: 0.9498 - fruit_output_loss: 0.1637 - loss: 0.1892 - val_fresh_output_accuracy: 0.9559 - val_fresh_output_loss: 0.1156 - val_fruit_output_accuracy: 0.9891 - val_fruit_output_loss: 0.0323 - val_loss: 0.0573 - learning_rate: 5.0000e-04

Epoch 4/8

3545/3545 ————— **179s** 50ms/step - fresh_output_accuracy: 0.9176
 - fresh_output_loss: 0.2020 - fruit_output_accuracy: 0.9711 - fruit_output_loss: 0.1078 - loss: 0.1361 - val_fresh_output_accuracy: 0.9592 - val_fresh_output_loss: 0.0861 - val_fruit_output_accuracy: 0.9921 - val_fruit_output_loss: 0.0259 - val_loss: 0.0440 - learning_rate: 5.0000e-04

Epoch 5/8

3545/3545 ————— **182s** 51ms/step - fresh_output_accuracy: 0.9263
 - fresh_output_loss: 0.1743 - fruit_output_accuracy: 0.9654 - fruit_output_loss:

ss: 0.1138 - loss: 0.1320 - val_fresh_output_accuracy: 0.9598 - val_fresh_output_loss: 0.0949 - val_fruit_output_accuracy: 0.9951 - val_fruit_output_loss: 0.0238 - val_loss: 0.0451 - learning_rate: 5.0000e-04

Epoch 6/8

3545/3545 ————— **181s** 51ms/step - fresh_output_accuracy: 0.9426 - fresh_output_loss: 0.1396 - fruit_output_accuracy: 0.9795 - fruit_output_loss: 0.0777 - loss: 0.0963 - val_fresh_output_accuracy: 0.9710 - val_fresh_output_loss: 0.0736 - val_fruit_output_accuracy: 0.9941 - val_fruit_output_loss: 0.0244 - val_loss: 0.0392 - learning_rate: 5.0000e-04

Epoch 7/8

3545/3545 ————— **181s** 51ms/step - fresh_output_accuracy: 0.9421 - fresh_output_loss: 0.1398 - fruit_output_accuracy: 0.9830 - fruit_output_loss: 0.0643 - loss: 0.0870 - val_fresh_output_accuracy: 0.9681 - val_fresh_output_loss: 0.0891 - val_fruit_output_accuracy: 0.9911 - val_fruit_output_loss: 0.0471 - val_loss: 0.0597 - learning_rate: 5.0000e-04

Epoch 8/8

3545/3545 ————— **180s** 51ms/step - fresh_output_accuracy: 0.9505 - fresh_output_loss: 0.1171 - fruit_output_accuracy: 0.9814 - fruit_output_loss: 0.0722 - loss: 0.0857 - val_fresh_output_accuracy: 0.9697 - val_fresh_output_loss: 0.0912 - val_fruit_output_accuracy: 0.9924 - val_fruit_output_loss: 0.0365 - val_loss: 0.0529 - learning_rate: 5.0000e-04

Restoring model weights from the end of the best epoch: 6.

Evaluating ResNet152...

✓ ResNet152 Results:
Test Fruit Acc: 0.9964
Test Fresh Acc: 0.9727

=====
Creating and training YOLOv3+FC

=====
✓ YOLOv3+FC created successfully!

Total parameters: 8,404,394

=====
Training YOLOv3+FC

=====
Epoch 1/8

3545/3545 ————— **171s** 39ms/step - fresh_output_accuracy: 0.5158 - fresh_output_loss: 0.7193 - fruit_output_accuracy: 0.1997 - fruit_output_loss: 2.0891 - loss: 2.4395 - val_fresh_output_accuracy: 0.5056 - val_fresh_output_loss: 0.6945 - val_fruit_output_accuracy: 0.2242 - val_fruit_output_loss: 4.5756 - val_loss: 3.4478 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **126s** 35ms/step - fresh_output_accuracy: 0.5231 - fresh_output_loss: 0.6940 - fruit_output_accuracy: 0.3076 - fruit_output_loss: 1.6585 - loss: 1.4038 - val_fresh_output_accuracy: 0.5540 - val_fresh_output_loss: 0.7158 - val_fruit_output_accuracy: 0.2406 - val_fruit_output_loss: 5.5920 - val_loss: 4.1700 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **125s** 35ms/step - fresh_output_accuracy: 0.5625 - fresh_output_loss: 0.6853 - fruit_output_accuracy: 0.3688 - fruit_output_loss: 1.4880 - loss: 1.2862 - val_fresh_output_accuracy: 0.5240 - val_fresh_output_loss: 0.6881 - val_fruit_output_accuracy: 0.3298 - val_fruit_output_loss: 1.6252 - val_loss: 1.3840 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **125s** 35ms/step - fresh_output_accuracy: 0.5716 - fresh_output_loss: 0.6807 - fruit_output_accuracy: 0.4083 - fruit_output_loss:

ss: 1.3736 - loss: 1.2067 - val_fresh_output_accuracy: 0.5128 - val_fresh_output_loss: 0.6930 - val_fruit_output_accuracy: 0.3647 - val_fruit_output_loss: 1.3530 - val_loss: 1.1969 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **125s** 35ms/step - fresh_output_accuracy: 0.5713 - fresh_output_loss: 0.6766 - fruit_output_accuracy: 0.4473 - fruit_output_loss: 1.2845 - loss: 1.1456 - val_fresh_output_accuracy: 0.5425 - val_fresh_output_loss: 0.6842 - val_fruit_output_accuracy: 0.3499 - val_fruit_output_loss: 2.2650 - val_loss: 1.8365 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **125s** 35ms/step - fresh_output_accuracy: 0.5760 - fresh_output_loss: 0.6720 - fruit_output_accuracy: 0.4878 - fruit_output_loss: 1.1966 - loss: 1.0868 - val_fresh_output_accuracy: 0.5859 - val_fresh_output_loss: 0.6662 - val_fruit_output_accuracy: 0.5797 - val_fruit_output_loss: 0.9747 - val_loss: 0.9317 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **126s** 35ms/step - fresh_output_accuracy: 0.5760 - fresh_output_loss: 0.6724 - fruit_output_accuracy: 0.5331 - fruit_output_loss: 1.1050 - loss: 1.0255 - val_fresh_output_accuracy: 0.5142 - val_fresh_output_loss: 0.7162 - val_fruit_output_accuracy: 0.3914 - val_fruit_output_loss: 1.4529 - val_loss: 1.2901 - learning_rate: 0.0010

Epoch 8/8

3545/3545 ————— **124s** 35ms/step - fresh_output_accuracy: 0.5734 - fresh_output_loss: 0.6713 - fruit_output_accuracy: 0.5720 - fruit_output_loss: 1.0340 - loss: 0.9800 - val_fresh_output_accuracy: 0.5204 - val_fresh_output_loss: 0.6941 - val_fruit_output_accuracy: 0.3690 - val_fruit_output_loss: 2.9198 - val_loss: 2.3100 - learning_rate: 0.0010

Restoring model weights from the end of the best epoch: 6.

Evaluating YOLOv3+FC...

✓ YOLOv3+FC Results:

Test Fruit Acc: 0.5851

Test Fresh Acc: 0.5986

=====
Creating and training EfficientNetB4

=====
Downloading data from https://storage.googleapis.com/keras-applications/efficientnetb4_notop.h5

71686520/71686520 ————— **0s** 0us/step

✓ EfficientNetB4 created successfully!

Total parameters: 18,775,785

=====
Training EfficientNetB4

=====
Epoch 1/8

3545/3545 ————— **199s** 40ms/step - fresh_output_accuracy: 0.5150 - fresh_output_loss: 0.7343 - fruit_output_accuracy: 0.2279 - fruit_output_loss: 1.9931 - loss: 1.6154 - val_fresh_output_accuracy: 0.5474 - val_fresh_output_loss: 0.6869 - val_fruit_output_accuracy: 0.3032 - val_fruit_output_loss: 1.6582 - val_loss: 1.3672 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **101s** 29ms/step - fresh_output_accuracy: 0.5505 - fresh_output_loss: 0.6873 - fruit_output_accuracy: 0.2919 - fruit_output_loss: 1.6936 - loss: 1.3917 - val_fresh_output_accuracy: 0.5484 - val_fresh_output_loss: 0.6835 - val_fruit_output_accuracy: 0.3752 - val_fruit_output_loss: 1.5186 - val_loss: 1.2681 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **102s** 29ms/step - fresh_output_accuracy: 0.5550
 - fresh_output_loss: 0.6832 - fruit_output_accuracy: 0.3066 - fruit_output_loss: 1.6374 - loss: 1.3511 - val_fresh_output_accuracy: 0.5688 - val_fresh_output_loss: 0.6697 - val_fruit_output_accuracy: 0.3275 - val_fruit_output_loss: 1.5163 - val_loss: 1.2623 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **102s** 29ms/step - fresh_output_accuracy: 0.5551
 - fresh_output_loss: 0.6755 - fruit_output_accuracy: 0.3353 - fruit_output_loss: 1.5810 - loss: 1.3094 - val_fresh_output_accuracy: 0.5234 - val_fresh_output_loss: 0.6831 - val_fruit_output_accuracy: 0.2462 - val_fruit_output_loss: 2.4584 - val_loss: 1.9256 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **103s** 29ms/step - fresh_output_accuracy: 0.5523
 - fresh_output_loss: 0.6753 - fruit_output_accuracy: 0.3355 - fruit_output_loss: 1.5555 - loss: 1.2914 - val_fresh_output_accuracy: 0.5500 - val_fresh_output_loss: 0.6863 - val_fruit_output_accuracy: 0.3937 - val_fruit_output_loss: 1.4061 - val_loss: 1.1901 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **106s** 30ms/step - fresh_output_accuracy: 0.5483
 - fresh_output_loss: 0.6740 - fruit_output_accuracy: 0.3490 - fruit_output_loss: 1.5140 - loss: 1.2619 - val_fresh_output_accuracy: 0.5718 - val_fresh_output_loss: 0.7261 - val_fruit_output_accuracy: 0.3989 - val_fruit_output_loss: 1.3977 - val_loss: 1.1963 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **103s** 29ms/step - fresh_output_accuracy: 0.5665
 - fresh_output_loss: 0.6640 - fruit_output_accuracy: 0.3495 - fruit_output_loss: 1.5096 - loss: 1.2559 - val_fresh_output_accuracy: 0.5606 - val_fresh_output_loss: 0.6581 - val_fruit_output_accuracy: 0.4490 - val_fruit_output_loss: 1.3732 - val_loss: 1.1586 - learning_rate: 0.0010

Epoch 8/8

3545/3545 ————— **104s** 29ms/step - fresh_output_accuracy: 0.5576
 - fresh_output_loss: 0.6700 - fruit_output_accuracy: 0.3547 - fruit_output_loss: 1.5091 - loss: 1.2574 - val_fresh_output_accuracy: 0.5754 - val_fresh_output_loss: 0.6933 - val_fruit_output_accuracy: 0.3726 - val_fruit_output_loss: 1.4202 - val_loss: 1.2021 - learning_rate: 0.0010

Restoring model weights from the end of the best epoch: 7.

Evaluating EfficientNetB4...

✓ EfficientNetB4 Results:

Test Fruit Acc: 0.4561

Test Fresh Acc: 0.5601

=====
 Creating and training DenseNet121

=====
 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/densenet/densenet121_weights_tf_dim_ordering_tf_kernels_notop.h5

29084464/29084464 ————— **0s** 0us/step

✓ DenseNet121 created successfully!

Total parameters: 7,746,250

=====
 Training DenseNet121

=====
 Epoch 1/8

3545/3545 ————— **166s** 36ms/step - fresh_output_accuracy: 0.7746
 - fresh_output_loss: 0.4628 - fruit_output_accuracy: 0.7938 - fruit_output_loss:

ss: 0.5989 - loss: 0.5578 - val_fresh_output_accuracy: 0.9068 - val_fresh_output_loss: 0.2250 - val_fruit_output_accuracy: 0.9740 - val_fruit_output_loss: 0.0736 - val_loss: 0.1191 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **98s** 28ms/step - fresh_output_accuracy: 0.8758 - fresh_output_loss: 0.2902 - fruit_output_accuracy: 0.9240 - fruit_output_loss: 0.2421 - loss: 0.2565 - val_fresh_output_accuracy: 0.9302 - val_fresh_output_loss: 0.1663 - val_fruit_output_accuracy: 0.9773 - val_fruit_output_loss: 0.0659 - val_loss: 0.0961 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **95s** 27ms/step - fresh_output_accuracy: 0.8886 - fresh_output_loss: 0.2446 - fruit_output_accuracy: 0.9448 - fruit_output_loss: 0.1785 - loss: 0.1984 - val_fresh_output_accuracy: 0.9319 - val_fresh_output_loss: 0.1952 - val_fruit_output_accuracy: 0.9872 - val_fruit_output_loss: 0.0427 - val_loss: 0.0885 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **95s** 27ms/step - fresh_output_accuracy: 0.9027 - fresh_output_loss: 0.2225 - fruit_output_accuracy: 0.9497 - fruit_output_loss: 0.1694 - loss: 0.1853 - val_fresh_output_accuracy: 0.9417 - val_fresh_output_loss: 0.1428 - val_fruit_output_accuracy: 0.9878 - val_fruit_output_loss: 0.0479 - val_loss: 0.0764 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **96s** 27ms/step - fresh_output_accuracy: 0.9134 - fresh_output_loss: 0.2002 - fruit_output_accuracy: 0.9634 - fruit_output_loss: 0.1280 - loss: 0.1496 - val_fresh_output_accuracy: 0.9477 - val_fresh_output_loss: 0.1123 - val_fruit_output_accuracy: 0.9914 - val_fruit_output_loss: 0.0355 - val_loss: 0.0586 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **95s** 27ms/step - fresh_output_accuracy: 0.9204 - fresh_output_loss: 0.1866 - fruit_output_accuracy: 0.9684 - fruit_output_loss: 0.1120 - loss: 0.1344 - val_fresh_output_accuracy: 0.9460 - val_fresh_output_loss: 0.1218 - val_fruit_output_accuracy: 0.9931 - val_fruit_output_loss: 0.0234 - val_loss: 0.0530 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **93s** 26ms/step - fresh_output_accuracy: 0.9222 - fresh_output_loss: 0.1866 - fruit_output_accuracy: 0.9645 - fruit_output_loss: 0.1276 - loss: 0.1453 - val_fresh_output_accuracy: 0.9562 - val_fresh_output_loss: 0.1072 - val_fruit_output_accuracy: 0.9924 - val_fruit_output_loss: 0.0272 - val_loss: 0.0512 - learning_rate: 0.0010

Epoch 8/8

3545/3545 ————— **95s** 27ms/step - fresh_output_accuracy: 0.9278 - fresh_output_loss: 0.1686 - fruit_output_accuracy: 0.9718 - fruit_output_loss: 0.1055 - loss: 0.1244 - val_fresh_output_accuracy: 0.9562 - val_fresh_output_loss: 0.1126 - val_fruit_output_accuracy: 0.9954 - val_fruit_output_loss: 0.0155 - val_loss: 0.0446 - learning_rate: 0.0010

Restoring model weights from the end of the best epoch: 8.

Evaluating DenseNet121...

✓ DenseNet121 Results:
Test Fruit Acc: 0.9944
Test Fresh Acc: 0.9566

=====
Creating and training ViT-Base Pretrained
=====

Error creating pretrained ViT model: https://tfhub.dev/google/vit_base_patch16_224/1 does not appear to be a valid module.

✓ ViT-Base Pretrained created successfully!

Total parameters: 27,048,746

=====
Training ViT-Base Pretrained
=====

Epoch 1/10

3545/3545 ————— **116s** 28ms/step - fresh_output_accuracy: 0.5388
- fresh_output_loss: 0.7226 - fruit_output_accuracy: 0.4163 - fruit_output_loss: 1.5209 - loss: 1.2815 - val_fresh_output_accuracy: 0.6346 - val_fresh_output_loss: 0.6364 - val_fruit_output_accuracy: 0.6698 - val_fruit_output_loss: 0.8535 - val_loss: 0.7883 - learning_rate: 5.0000e-04

Epoch 2/10

3545/3545 ————— **92s** 26ms/step - fresh_output_accuracy: 0.6243
- fresh_output_loss: 0.6415 - fruit_output_accuracy: 0.5965 - fruit_output_loss: 1.0313 - loss: 0.9143 - val_fresh_output_accuracy: 0.6406 - val_fresh_output_loss: 1.4673 - val_fruit_output_accuracy: 0.6705 - val_fruit_output_loss: 1.1979 - val_loss: 1.2790 - learning_rate: 5.0000e-04

Epoch 3/10

3545/3545 ————— **89s** 25ms/step - fresh_output_accuracy: 0.6331
- fresh_output_loss: 0.6209 - fruit_output_accuracy: 0.6211 - fruit_output_loss: 0.9417 - loss: 0.8455 - val_fresh_output_accuracy: 0.6303 - val_fresh_output_loss: 1.6012 - val_fruit_output_accuracy: 0.6666 - val_fruit_output_loss: 12.1641 - val_loss: 9.0007 - learning_rate: 5.0000e-04

Epoch 4/10

3545/3545 ————— **88s** 25ms/step - fresh_output_accuracy: 0.6472
- fresh_output_loss: 0.6058 - fruit_output_accuracy: 0.6518 - fruit_output_loss: 0.8429 - loss: 0.7715 - val_fresh_output_accuracy: 0.7222 - val_fresh_output_loss: 0.5317 - val_fruit_output_accuracy: 0.7344 - val_fruit_output_loss: 0.5947 - val_loss: 0.5760 - learning_rate: 5.0000e-04

Epoch 5/10

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.6740
- fresh_output_loss: 0.5840 - fruit_output_accuracy: 0.6775 - fruit_output_loss: 0.7827 - loss: 0.7231 - val_fresh_output_accuracy: 0.7732 - val_fresh_output_loss: 0.4601 - val_fruit_output_accuracy: 0.7857 - val_fruit_output_loss: 0.4904 - val_loss: 0.4814 - learning_rate: 5.0000e-04

Epoch 6/10

3545/3545 ————— **90s** 25ms/step - fresh_output_accuracy: 0.7160
- fresh_output_loss: 0.5373 - fruit_output_accuracy: 0.7075 - fruit_output_loss: 0.7295 - loss: 0.6719 - val_fresh_output_accuracy: 0.7801 - val_fresh_output_loss: 0.4137 - val_fruit_output_accuracy: 0.8223 - val_fruit_output_loss: 0.5090 - val_loss: 0.4807 - learning_rate: 5.0000e-04

Epoch 7/10

3545/3545 ————— **89s** 25ms/step - fresh_output_accuracy: 0.7104
- fresh_output_loss: 0.5341 - fruit_output_accuracy: 0.7052 - fruit_output_loss: 0.7360 - loss: 0.6754 - val_fresh_output_accuracy: 0.6916 - val_fresh_output_loss: 0.5785 - val_fruit_output_accuracy: 0.7255 - val_fruit_output_loss: 0.6866 - val_loss: 0.6542 - learning_rate: 5.0000e-04

Epoch 8/10

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.6345
- fresh_output_loss: 0.6080 - fruit_output_accuracy: 0.6377 - fruit_output_loss: 0.8627 - loss: 0.7863 - val_fresh_output_accuracy: 0.7383 - val_fresh_output_loss: 0.4954 - val_fruit_output_accuracy: 0.7683 - val_fruit_output_loss: 0.5342 - val_loss: 0.5226 - learning_rate: 5.0000e-04

Epoch 9/10

3545/3545 ————— **90s** 25ms/step - fresh_output_accuracy: 0.7054
- fresh_output_loss: 0.5432 - fruit_output_accuracy: 0.7022 - fruit_output_loss: 0.7116 - loss: 0.6611 - val_fresh_output_accuracy: 0.7959 - val_fresh_out

put_loss: 0.4312 - val_fruit_output_accuracy: 0.8331 - val_fruit_output_loss: 0.4133 - val_loss: 0.4188 - learning_rate: 5.0000e-04

Epoch 10/10

3545/3545 ————— **91s** 26ms/step - fresh_output_accuracy: 0.7396 - fresh_output_loss: 0.5012 - fruit_output_accuracy: 0.7556 - fruit_output_loss: 0.6005 - loss: 0.5707 - val_fresh_output_accuracy: 0.7828 - val_fresh_output_loss: 0.4129 - val_fruit_output_accuracy: 0.8371 - val_fruit_output_loss: 0.5322 - val_loss: 0.4966 - learning_rate: 5.0000e-04

Restoring model weights from the end of the best epoch: 9.

Evaluating ViT-Base Pretrained...

✓ ViT-Base Pretrained Results:

Test Fruit Acc: 0.8414

Test Fresh Acc: 0.7881

=====
Creating and training ConvNeXt-Base

=====

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/convnext/convnext_base_notop.h5

350926856/350926856 ————— **2s** 0us/step

✓ ConvNeXt-Base created successfully!

Total parameters: 88,275,210

=====
Training ConvNeXt-Base

=====

Epoch 1/8

E0000 00:00:1754730424.493469 106 gpu_timer.cc:82] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.

E0000 00:00:1754730424.678580 106 gpu_timer.cc:82] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.

3545/3545 ————— **248s** 59ms/step - fresh_output_accuracy: 0.6775
 - fresh_output_loss: 0.5955 - fruit_output_accuracy: 0.7018 - fruit_output_loss: 0.8983 - loss: 0.8074 - val_fresh_output_accuracy: 0.8009 - val_fresh_output_loss: 0.3926 - val_fruit_output_accuracy: 0.8927 - val_fruit_output_loss: 0.2878 - val_loss: 0.3194 - learning_rate: 0.0010

Epoch 2/8

3545/3545 ————— **195s** 55ms/step - fresh_output_accuracy: 0.7854
 - fresh_output_loss: 0.4415 - fruit_output_accuracy: 0.8679 - fruit_output_loss: 0.4105 - loss: 0.4198 - val_fresh_output_accuracy: 0.8288 - val_fresh_output_loss: 0.3609 - val_fruit_output_accuracy: 0.9296 - val_fruit_output_loss: 0.2087 - val_loss: 0.2545 - learning_rate: 0.0010

Epoch 3/8

3545/3545 ————— **195s** 55ms/step - fresh_output_accuracy: 0.8150
 - fresh_output_loss: 0.3950 - fruit_output_accuracy: 0.8932 - fruit_output_loss: 0.3392 - loss: 0.3560 - val_fresh_output_accuracy: 0.8746 - val_fresh_output_loss: 0.2699 - val_fruit_output_accuracy: 0.9286 - val_fruit_output_loss: 0.2616 - val_loss: 0.2643 - learning_rate: 0.0010

Epoch 4/8

3545/3545 ————— **195s** 55ms/step - fresh_output_accuracy: 0.8540
 - fresh_output_loss: 0.3363 - fruit_output_accuracy: 0.9126 - fruit_output_loss: 0.2819 - loss: 0.2983 - val_fresh_output_accuracy: 0.8910 - val_fresh_output_loss: 0.2423 - val_fruit_output_accuracy: 0.9447 - val_fruit_output_loss: 0.1784 - val_loss: 0.1977 - learning_rate: 0.0010

Epoch 5/8

3545/3545 ————— **195s** 55ms/step - fresh_output_accuracy: 0.8614
 - fresh_output_loss: 0.3164 - fruit_output_accuracy: 0.9167 - fruit_output_loss: 0.2734 - loss: 0.2864 - val_fresh_output_accuracy: 0.8864 - val_fresh_output_loss: 0.2542 - val_fruit_output_accuracy: 0.9361 - val_fruit_output_loss: 0.2309 - val_loss: 0.2371 - learning_rate: 0.0010

Epoch 6/8

3545/3545 ————— **193s** 55ms/step - fresh_output_accuracy: 0.8585
 - fresh_output_loss: 0.3102 - fruit_output_accuracy: 0.9198 - fruit_output_loss: 0.2732 - loss: 0.2842 - val_fresh_output_accuracy: 0.9029 - val_fresh_output_loss: 0.2217 - val_fruit_output_accuracy: 0.9394 - val_fruit_output_loss: 0.2096 - val_loss: 0.2134 - learning_rate: 0.0010

Epoch 7/8

3545/3545 ————— **0s** 48ms/step - fresh_output_accuracy: 0.8761 - fresh_output_loss: 0.2770 - fruit_output_accuracy: 0.9295 - fruit_output_loss: 0.2304 - loss: 0.2444

Epoch 7: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

3545/3545 ————— **195s** 55ms/step - fresh_output_accuracy: 0.8761
 - fresh_output_loss: 0.2770 - fruit_output_accuracy: 0.9295 - fruit_output_loss: 0.2304 - loss: 0.2444 - val_fresh_output_accuracy: 0.8772 - val_fresh_output_loss: 0.3073 - val_fruit_output_accuracy: 0.8733 - val_fruit_output_loss: 0.3856 - val_loss: 0.3596 - learning_rate: 0.0010

Epoch 8/8

3545/3545 ————— **197s** 56ms/step - fresh_output_accuracy: 0.8814
 - fresh_output_loss: 0.2661 - fruit_output_accuracy: 0.9430 - fruit_output_loss: 0.1899 - loss: 0.2128 - val_fresh_output_accuracy: 0.9144 - val_fresh_output_loss: 0.2005 - val_fruit_output_accuracy: 0.9700 - val_fruit_output_loss: 0.1229 - val_loss: 0.1463 - learning_rate: 5.0000e-04

Restoring model weights from the end of the best epoch: 8.

Evaluating ConvNeXt-Base...

✓ ConvNeXt-Base Results:

Test Fruit Acc: 0.9691

Test Fresh Acc: 0.9184

```
=====
===
🏆 BEST MODEL: ResNet152
Test Fruit Accuracy: 0.9964
Test Fresh Accuracy: 0.9727
Combined Test Accuracy: 0.9845
=====
===
```


```
=====
FINAL MULTI-TASK LEARNING SUMMARY TABLE
=====
```

	Model	Train Loss	Valid Loss	Test Loss	Train Fruit Acc	Valid Fruit Acc	Test Fruit Acc	Train Fresh Acc	Valid Fresh Acc	Test Fresh Acc	Combined
	Custom CNN	0.6229	0.6403	0.6280	0.7110	0.7154	0.6873	0.6764	0.6808	0.6981	
	VGG16	0.4495	0.3433	0.3308	0.9236	0.9296	0.5665	0.5642	0.5676	0.7486	
	InceptionV3	0.0316	0.0551	0.0495	0.9934	0.9908	0.9599	0.9510	0.9552	0.9730	
	ResNet152	0.0190	0.0392	0.0282	0.9941	0.9964	0.9747	0.9710	0.9727	0.9845	
	YOLOv3+FC	0.9220	0.9317	0.9275	0.5797	0.5851	0.5889	0.5859	0.5814	0.5918	
	EfficientNetB4	1.1444	1.1586	1.1351	0.4490	0.4561	0.5655	0.5606	0.4508	0.5081	
	DenseNet121	0.0308	0.0446	0.0400	0.9954	0.9944	0.9606	0.9562	0.9566	0.9755	
	ViT-Base Pretrained	0.5120	0.4188	0.4447	0.8331	0.8414	0.7986	0.7959	0.8414	0.8147	
	ConvNeXt-Base	0.0903	0.1463	0.1415	0.9700	0.9691	0.9285	0.9144	0.9800	0.9437	


✓ Results saved to 'multi_task_model_comparison_results.csv'

```
=====
CREATING CONFERENCE-READY VISUALIZATIONS
=====
```


Created output directory: conference_figures




✓ Fruit accuracy comparison saved: conference_figures/fruit_classification_accuracy_comparison.png




✓ Freshness accuracy comparison saved: conference_figures/freshness_classification_accuracy_comparison.png




✓ Combined performance comparison saved: conference_figures/combined_performance_comparison.png



✓ Loss comparison saved: conference_figures/loss_comparison.png




✓ Fruit classification heatmap saved: conference_figures/fruit_classification_heatmap.png




✓ Freshness classification heatmap saved: conference_figures/freshness_classification_heatmap.png

Multi-Task Performance Correlation Analysis



✓ Performance correlation scatter saved: conference_figures/performance_correlation_scatter.png


Average Performance Across All Dataset Splits



✓ Average performance comparison saved: conference_figures/average_performance_comparison.png




✓ Training history saved: conference_figures/custom_cnn_training_history.png

VGG16 - Training History Analysis


✓ Training history saved: conference_figures/vgg16_training_history.png




✓ Training history saved: conference_figures/inceptionv3_training_history.png

ResNet152 - Training History Analysis


✓ Training history saved: conference_figures/resnet152_training_history.png

YOLOv3+FC - Training History Analysis


✓ Training history saved: conference_figures/yolov3_fc_training_history.png

EfficientNetB4 - Training History Analysis


✓ Training history saved: conference_figures/efficientnetb4_training_history.png




✓ Training history saved: conference_figures/densenet121_training_history.png




✓ Training history saved: conference_figures/vit_base_pretrained_training_history.png

ConvNeXt-Base - Training History Analysis

✓ Training history saved: conference_figures/convnext_base_training_history.png


Best Model (ResNet152) - Training History Analysis

✓ Best model training history saved: conference_figures/best_training_history.png
Creating confusion matrices for best model (ResNet152)...




✓ Best model fruit confusion matrix saved: conference_figures/best_fruit_confusion_matrix.png

Best Model (ResNet152) - Freshness Classification Confusion Matrix




✓ Best model freshness confusion matrix saved: conference_figures/best_freshness_confusion_matrix.png
Creating ROC curve for best model (ResNet152)...


Best Model (ResNet152) - Freshness Classification ROC Curve




✓ Best model ROC curve saved: conference_figures/best_freshness_roc_curve.png
Creating classification reports for best model (ResNet152)...




✓ Best model classification reports saved: conference_figures/best_classification_reports.png
Creating confidence distribution for best model (ResNet152)...



✓ Best model confidence distribution saved: conference_figures/best_prediction_confidence_distribution.png
Creating sample predictions for best model (ResNet152)...


Best Model (ResNet152) - Sample Predictions

✓ Best model sample predictions saved: conference_figures/best_sample_predictions.png
Creating confusion matrices for ViT-Base Pretrained...




✓ ViT-Base Pretrained fruit confusion matrix saved: conference_figures/vit_base_pretrained_fruit_confusion_matrix.png

ViT-Base Pretrained - Freshness Classification Confusion Matrix



✓ ViT-Base Pretrained freshness confusion matrix saved: conference_figures/vi_t_base_pretrained_freshness_confusion_matrix.png



✓ Results summary visualization saved: conference_figures/multi_task_results_summary.png

=====

ALL VISUALIZATIONS SAVED SUCCESSFULLY!

Check the 'conference_figures' folder for all conference-ready figures

=====

=====

===

✅ ENHANCED MULTI-TASK LEARNING PROJECT COMPLETED!

🏆 Best Model: ResNet152

🎯 Target fruit classes: ['apple', 'banana', 'bittergourd', 'capsicum', 'cucumber', 'okra', 'orange', 'potato', 'tomato']

📊 Total models trained: 9

📁 All visualizations saved in 'conference_figures' folder

=====

===

=====

===

🎉 PROJECT COMPLETED SUCCESSFULLY! 🎉

=====

===

📊 Results available in DataFrame with 9 models

🏆 Best performing model: ResNet152

📈 Best combined accuracy: 0.9845

🍏 Best fruit accuracy: 0.9964

🌱 Best fresh accuracy: 0.9727

🎯 Fruit classes handled: 9

📋 DETAILED RESULTS:

- =====
- Custom CNN:
 - ↳ Fruit: 0.7154, Fresh: 0.6808, Combined: 0.6981
 - VGG16:
 - ↳ Fruit: 0.9296, Fresh: 0.5676, Combined: 0.7486
 - InceptionV3:
 - ↳ Fruit: 0.9908, Fresh: 0.9552, Combined: 0.9730
 - ResNet152:
 - ↳ Fruit: 0.9964, Fresh: 0.9727, Combined: 0.9845
 - YOLOv3+FC:
 - ↳ Fruit: 0.5851, Fresh: 0.5986, Combined: 0.5918
 - EfficientNetB4:
 - ↳ Fruit: 0.4561, Fresh: 0.5601, Combined: 0.5081
 - DenseNet121:
 - ↳ Fruit: 0.9944, Fresh: 0.9566, Combined: 0.9755
 - ViT-Base Pretrained:
 - ↳ Fruit: 0.8414, Fresh: 0.7881, Combined: 0.8147
 - ConvNeXt-Base:
 - ↳ Fruit: 0.9691, Fresh: 0.9184, Combined: 0.9437

📁 FILES GENERATED:

- =====
- multi_task_model_comparison_results.csv - Main results
 - conference_figures/ - All visualization files

 VISUALIZATIONS CREATED:

=====

- Individual accuracy comparisons
- Combined performance plots
- Loss comparisons
- Performance heatmaps
- Correlation scatter plots
- Average performance bars
- Training history for each model
- Best model training history
- Best model confusion matrices
- Best model ROC curves
- Best model classification reports
- Best model confidence distributions
- Best model sample predictions
- ViT-Base Pretrained confusion matrices
- Comprehensive results summary

=====

===

=====

===

 ENHANCED MULTI-TASK FRUIT CLASSIFICATION PROJECT WITH ALL ISSUES FIXED FOR 9 MODELS - COMPLETED!

=====

===

```

In [2]: # =====
# SAVE ALL TRAINED MODELS AS .KERAS FILES
# Run this code in the SECOND BLOCK after training all models
# =====

import os
import tensorflow as tf

# Create save directory
save_dir = "/kaggle/working/"
os.makedirs(save_dir, exist_ok=True)

print("="*70)
print("📁 SAVING ALL TRAINED MODELS TO .KERAS FILES")
print("="*70)

# Check if we have trained models from the previous block
try:
    if 'trained_models' not in locals() and 'trained_models' not in globals():
        print("❌ Error: trained_models not found. Make sure you ran the training code first.")
    else:
        # Save all trained models
        saved_models_info = []
        failed_models = []

        for model_name, model in trained_models.items():
            try:
                # Create safe filename (remove spaces and special characters)
                safe_name = model_name.replace(' ', '_').replace('+', '_').replace('-', '_').replace('/', '_')
                save_path = os.path.join(save_dir, f"{safe_name}.keras")

                # Save the model in .keras format
                model.save(save_path)

                # Get file size for confirmation
                file_size = os.path.getsize(save_path) / (1024 * 1024) # Size in MB

                saved_models_info.append({
                    'name': model_name,
                    'safe_name': safe_name,
                    'path': save_path,
                    'size_mb': file_size
                })

                print(f"✅ {model_name:20} → {safe_name}.keras ({file_size:.1f} MB)")

            except Exception as e:
                failed_models.append({
                    'name': model_name,
                    'error': str(e)
                })
                print(f"❌ {model_name:20} → Failed: {str(e)}")

```

```

# Save the best model with special naming (if available)
if 'best_model' in locals() or 'best_model' in globals():
    if best_model is not None and 'best_model_name' in locals():
        try:
            best_safe_name = best_model_name.replace(' ', '_').replace(
                '+', '_').replace('-', '_').replace('/', '_')
            best_save_path = os.path.join(save_dir, f"BEST_{best_safe_
name}.keras")
            best_model.save(best_save_path)

            best_file_size = os.path.getsize(best_save_path) / (1024 *
1024)
            saved_models_info.append({
                'name': f"BEST_{best_model_name}",
                'safe_name': f"BEST_{best_safe_name}",
                'path': best_save_path,
                'size_mb': best_file_size
            })
            print(f"🏆 BEST MODEL: {best_model_name:12} → BEST_{best_
safe_name}.keras ({best_file_size:.1f} MB)")
        except Exception as e:
            print(f"❌ Failed to save best model: {str(e)}")

# Print summary
print("\n" + "="*70)
print("📁 MODEL SAVING SUMMARY")
print("="*70)

if saved_models_info:
    print(f"✅ Successfully saved: {len(saved_models_info)} models")
    print(f"📁 Save location: {save_dir}")
    print(f"📦 Total size: {sum([info['size_mb'] for info in saved_mo
dels_info]):.1f} MB")

    print(f"\n📄 SAVED FILES:")
    print("-" * 50)
    for info in saved_models_info:
        print(f" • {info['safe_name']}.keras ({info['size_mb']:.1f} M
B)")

if failed_models:
    print(f"\n❌ Failed to save: {len(failed_models)} models")
    for failed in failed_models:
        print(f" • {failed['name']}: {failed['error']}")

# Show how to Load models Later
print(f"\n📄 TO LOAD MODELS LATER, USE:")
print("-" * 40)
print("import tensorflow as tf")
print("# Load any specific model:")
if saved_models_info:
    example_model = saved_models_info[0]
    print(f"model = tf.keras.models.load_model('{example_model['pat
h']}')")
print("# Or load the best model:")
if any('BEST_' in info['safe_name'] for info in saved_models_info):

```

```

        best_info = next(info for info in saved_models_info if 'BEST_' in
info['safe_name'])
        print(f"best_model = tf.keras.models.load_model('{best_info['pat
h']}')")

        print("\n" + "="*70)
        print("🎉 MODEL SAVING COMPLETED!")
        print("="*70)

    except Exception as e:
        print(f"❌ Critical error during model saving: {str(e)}")
        print("Make sure you have run the training code in the previous block firs
t.")

# Optional: List all files in the working directory to confirm
print(f"\n📁 ALL FILES IN {save_dir}:")
try:
    files = [f for f in os.listdir(save_dir) if f.endswith('.keras')]
    if files:
        for i, file in enumerate(files, 1):
            file_path = os.path.join(save_dir, file)
            file_size = os.path.getsize(file_path) / (1024 * 1024)
            print(f" {i:2d}. {file} ({file_size:.1f} MB)")
    else:
        print(" No .keras files found.")
except Exception as e:
    print(f" Error listing files: {e}")

```

SAVING ALL TRAINED MODELS TO .KERAS FILES

✓ Custom CNN	→ Custom_CNN.keras (309.7 MB)
✓ VGG16	→ VGG16.keras (115.4 MB)
✓ InceptionV3	→ InceptionV3.keras (113.1 MB)
✓ ResNet152	→ ResNet152.keras (348.8 MB)
✓ YOLOv3+FC	→ YOLOv3_FC.keras (96.5 MB)
✓ EfficientNetB4	→ EfficientNetB4.keras (116.6 MB)
✓ DenseNet121	→ DenseNet121.keras (38.9 MB)
✓ ViT-Base Pretrained	→ ViT_Base_Pretrained.keras (309.7 MB)
✓ ConvNeXt-Base	→ ConvNeXt_Base.keras (407.6 MB)

MODEL SAVING SUMMARY

- ✓ Successfully saved: 9 models
- 📁 Save location: /kaggle/working/
- 📁 Total size: 1856.3 MB

SAVED FILES:

- Custom_CNN.keras (309.7 MB)
- VGG16.keras (115.4 MB)
- InceptionV3.keras (113.1 MB)
- ResNet152.keras (348.8 MB)
- YOLOv3_FC.keras (96.5 MB)
- EfficientNetB4.keras (116.6 MB)
- DenseNet121.keras (38.9 MB)
- ViT_Base_Pretrained.keras (309.7 MB)
- ConvNeXt_Base.keras (407.6 MB)

TO LOAD MODELS LATER, USE:

```
import tensorflow as tf
# Load any specific model:
model = tf.keras.models.load_model('/kaggle/working/Custom_CNN.keras')
# Or load the best model:
```

MODEL SAVING COMPLETED!

- 📁 ALL FILES IN /kaggle/working/:
 1. InceptionV3.keras (113.1 MB)
 2. VGG16.keras (115.4 MB)
 3. Custom_CNN.keras (309.7 MB)
 4. ViT_Base_Pretrained.keras (309.7 MB)
 5. ConvNeXt_Base.keras (407.6 MB)
 6. EfficientNetB4.keras (116.6 MB)
 7. ResNet152.keras (348.8 MB)
 8. DenseNet121.keras (38.9 MB)
 9. YOLOv3_FC.keras (96.5 MB)

In []: